

PyTorch

Tensor Algebra

Tensors are **multi-dimensional arrays** of numbers that generalize scalars, vectors, and matrices to higher dimensions.

Tensor algebra is a **branch** of **mathematics** that deals with **tensors** and their mathematical **operations**.

(11)

5	3	7
---	---	---

Scalar

5
1.5
2

Column Vector
(shape 3×1)

4	19	8
16	3	5

Matrix

A 3D cube is shown with faces colored purple, pink, and orange. The purple faces are numbered 1 through 9. The pink faces are labeled with letters a, b, c, c, g, j. The orange faces are labeled with letters A, B, C, C, G, J.

Tensor

Tensor Algebra: Addition and multiplication with scalars

To perform an **addition** or **multiplication** between a **scalar** and a **tensor** simply add or multiply each value by it

$$a + \begin{pmatrix} m_{00} & m_{01} & m_{02} \\ m_{10} & m_{11} & m_{12} \\ m_{20} & m_{21} & m_{22} \end{pmatrix} = \begin{pmatrix} a + m_{00} & a + m_{01} & a + m_{02} \\ a + m_{10} & a + m_{11} & a + m_{12} \\ a + m_{20} & a + m_{21} & a + m_{22} \end{pmatrix}$$

$$a \cdot \begin{pmatrix} m_{00} & m_{01} & m_{02} \\ m_{10} & m_{11} & m_{12} \\ m_{20} & m_{21} & m_{22} \end{pmatrix} = \begin{pmatrix} a \cdot m_{00} & a \cdot m_{01} & a \cdot m_{02} \\ a \cdot m_{10} & a \cdot m_{11} & a \cdot m_{12} \\ a \cdot m_{20} & a \cdot m_{21} & a \cdot m_{22} \end{pmatrix}$$

Tensor Algebra: Addition

To **add** two **tensor** simply **add** each **respective element**. For $C = A + B$:

$$C_{ij} = A_{ij} + B_{ij}$$

$$\begin{pmatrix} m_{00} & m_{01} & m_{02} \\ m_{10} & m_{11} & m_{12} \\ m_{20} & m_{21} & m_{22} \end{pmatrix} + \begin{pmatrix} h_{00} & h_{01} & h_{02} \\ h_{10} & h_{11} & h_{12} \\ h_{20} & h_{21} & h_{22} \end{pmatrix} = \begin{pmatrix} m_{00} + h_{00} & m_{01} + h_{01} & m_{02} + h_{02} \\ m_{10} + h_{10} & m_{11} + h_{11} & m_{12} + h_{12} \\ m_{20} + h_{20} & m_{21} + h_{21} & m_{22} + h_{22} \end{pmatrix}$$

Only tensors with the same shape can be added

Tensor Algebra: Broadcasting

Only tensors with the same shape can be added? that is **not 100% true**.
Different tensors can be added exploiting the **broadcast** operation.

Broadcasting is an implicit operation where a **tensor** is **automatically replicated** to **match** the **dimension** of the other operand **tensor**.

```
import torch

a = torch.tensor([[1,2,3],[4,5,6]])
b = torch.tensor([[1,2,3]])

print("shape a =", a.shape)
print("shape b =", b.shape)

c = a + b
print("c =", c)
print("shape c =", c.shape)
```

```
shape a = torch.Size([2, 3])
shape b = torch.Size([1, 3])
c = tensor([[2, 4, 6],
           [5, 7, 9]])
shape c = torch.Size([2, 3])
```

```
import torch

a = torch.tensor([[1,2,3],[4,5,6]])
b = torch.tensor([[1,2,3], [1,2,3]])

print("shape a =", a.shape)
print("shape b =", b.shape)

c = a + b
print("c =", c)
print("shape c =", c.shape)
```

```
shape a = torch.Size([2, 3])
shape b = torch.Size([2, 3])
c = tensor([[2, 4, 6],
           [5, 7, 9]])
shape c = torch.Size([2, 3])
```

Tensor Algebra: Multiplication

To **multiply** two **tensor** simply **multiply** each **respective element**. For $C = A \cdot B$:

$$C_{ij} = A_{ij} \cdot B_{ij}$$

$$\begin{pmatrix} m_{00} & m_{01} & m_{02} \\ m_{10} & m_{11} & m_{12} \\ m_{20} & m_{21} & m_{22} \end{pmatrix} \cdot \begin{pmatrix} h_{00} & h_{01} & h_{02} \\ h_{10} & h_{11} & h_{12} \\ h_{20} & h_{21} & h_{22} \end{pmatrix} = \begin{pmatrix} m_{00} \cdot h_{00} & m_{01} \cdot h_{01} & m_{02} \cdot h_{02} \\ m_{10} \cdot h_{10} & m_{11} \cdot h_{11} & m_{12} \cdot h_{12} \\ m_{20} \cdot h_{20} & m_{21} \cdot h_{21} & m_{22} \cdot h_{22} \end{pmatrix}$$

Tensor Algebra: Multiplication with broadcasting

The broadcast operator also works for multiplication

```
import torch

a = torch.tensor([[1,2,3],[4,5,6]])
b = torch.tensor([[1,2,3]])

print("shape a =", a.shape)
print("shape b =", b.shape)

c = a * b
print("c =", c)
print("shape c =", c.shape)
```

```
↳ shape a = torch.Size([2, 3])
   shape b = torch.Size([1, 3])
   c = tensor([[ 1,  4,  9],
               [ 4, 10, 18]])
   shape c = torch.Size([2, 3])
```

```
import torch

a = torch.tensor([[1,2,3],[4,5,6]])
b = torch.tensor([[1,2,3],[1,2,3]])

print("shape a =", a.shape)
print("shape b =", b.shape)

c = a * b
print("c =", c)
print("shape c =", c.shape)
```

```
↳ shape a = torch.Size([2, 3])
   shape b = torch.Size([2, 3])
   c = tensor([[ 1,  4,  9],
               [ 4, 10, 18]])
   shape c = torch.Size([2, 3])
```

Tensor Algebra: Multiplication with broadcasting (2)

```
▶ import torch

a = torch.tensor([[1,2,3]])

print("shape a =", a.shape)
print("shape a.T =", a.T.shape)

c = a * a.T
print("c =", c)
print("shape c =", c.shape)
```

```
↳ shape a = torch.Size([1, 3])
   shape a.T = torch.Size([3, 1])
   c = tensor([[1, 2, 3],
               [2, 4, 6],
               [3, 6, 9]])
   shape c = torch.Size([3, 3])
```

```
▶ import torch

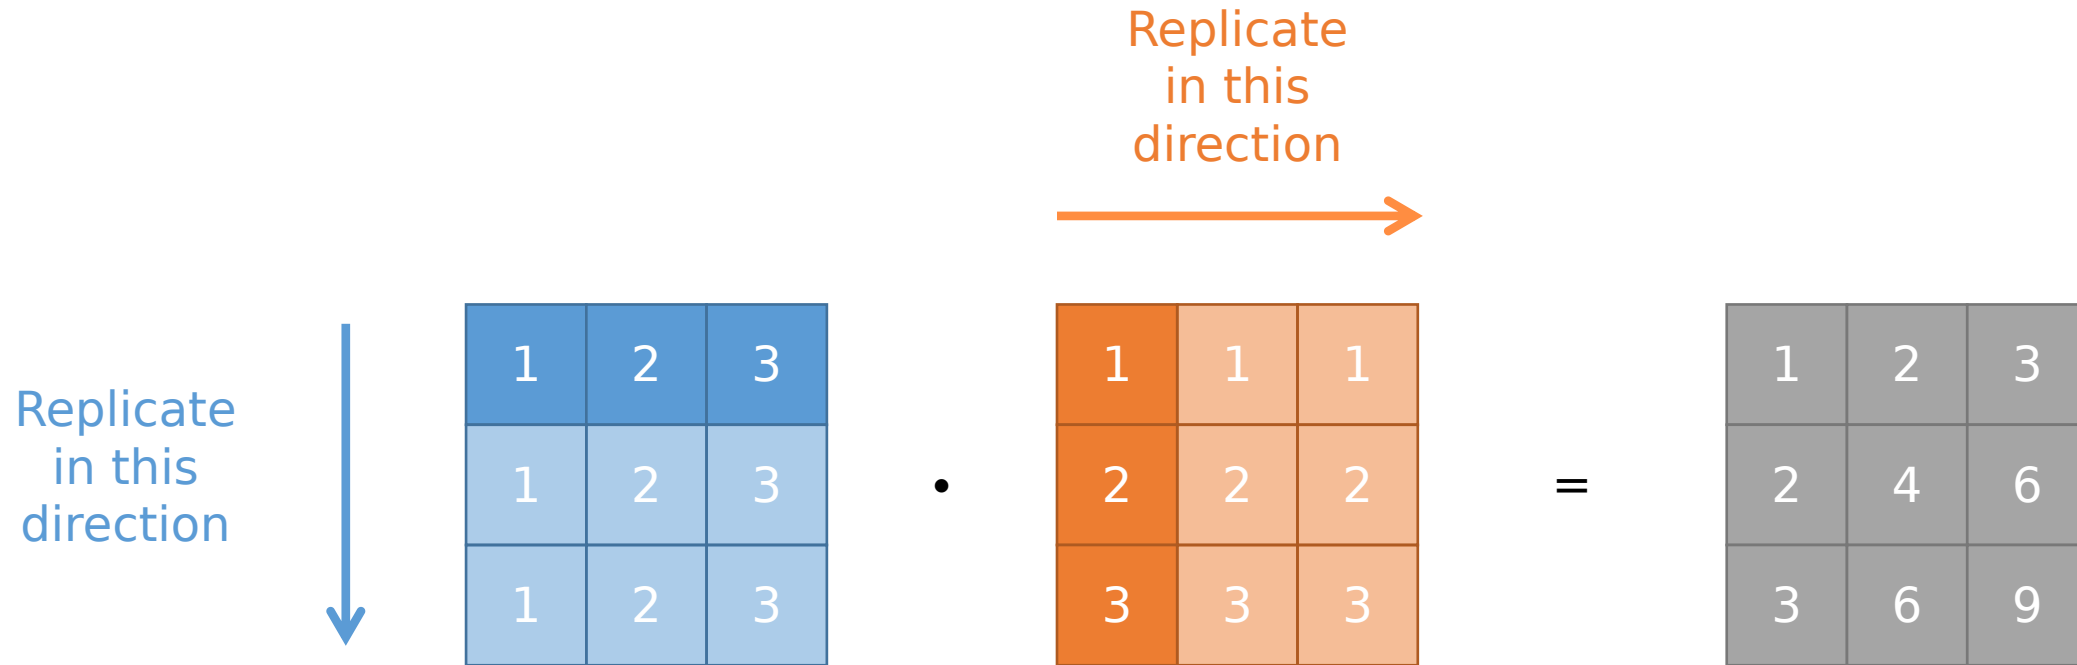
a = torch.tensor([[1,2,3], [1,2,3], [1,2,3]])
b = torch.tensor([[1,1,1], [2,2,2], [3,3,3]])

print("shape a =", a.shape)
print("shape b =", b.shape)

c = a * b
print("c =", c)
print("shape c =", c.shape)
```

```
↳ shape a = torch.Size([3, 3])
   shape b = torch.Size([3, 3])
   c = tensor([[1, 2, 3],
               [2, 4, 6],
               [3, 6, 9]])
   shape c = torch.Size([3, 3])
```


Tensor Algebra: Multiplication with broadcasting (3)



Tensor Algebra: Rearrange

Manipulating the **shape** of a **tensor** is called a rearrange operation.
When **rearranging** a tensor the **number of elements** remains the **same**.
The **product** of the **dimensions** is then **constant**.

1	2	3	4
5	6	7	8
9	10	11	12

Shape = (3, 4)

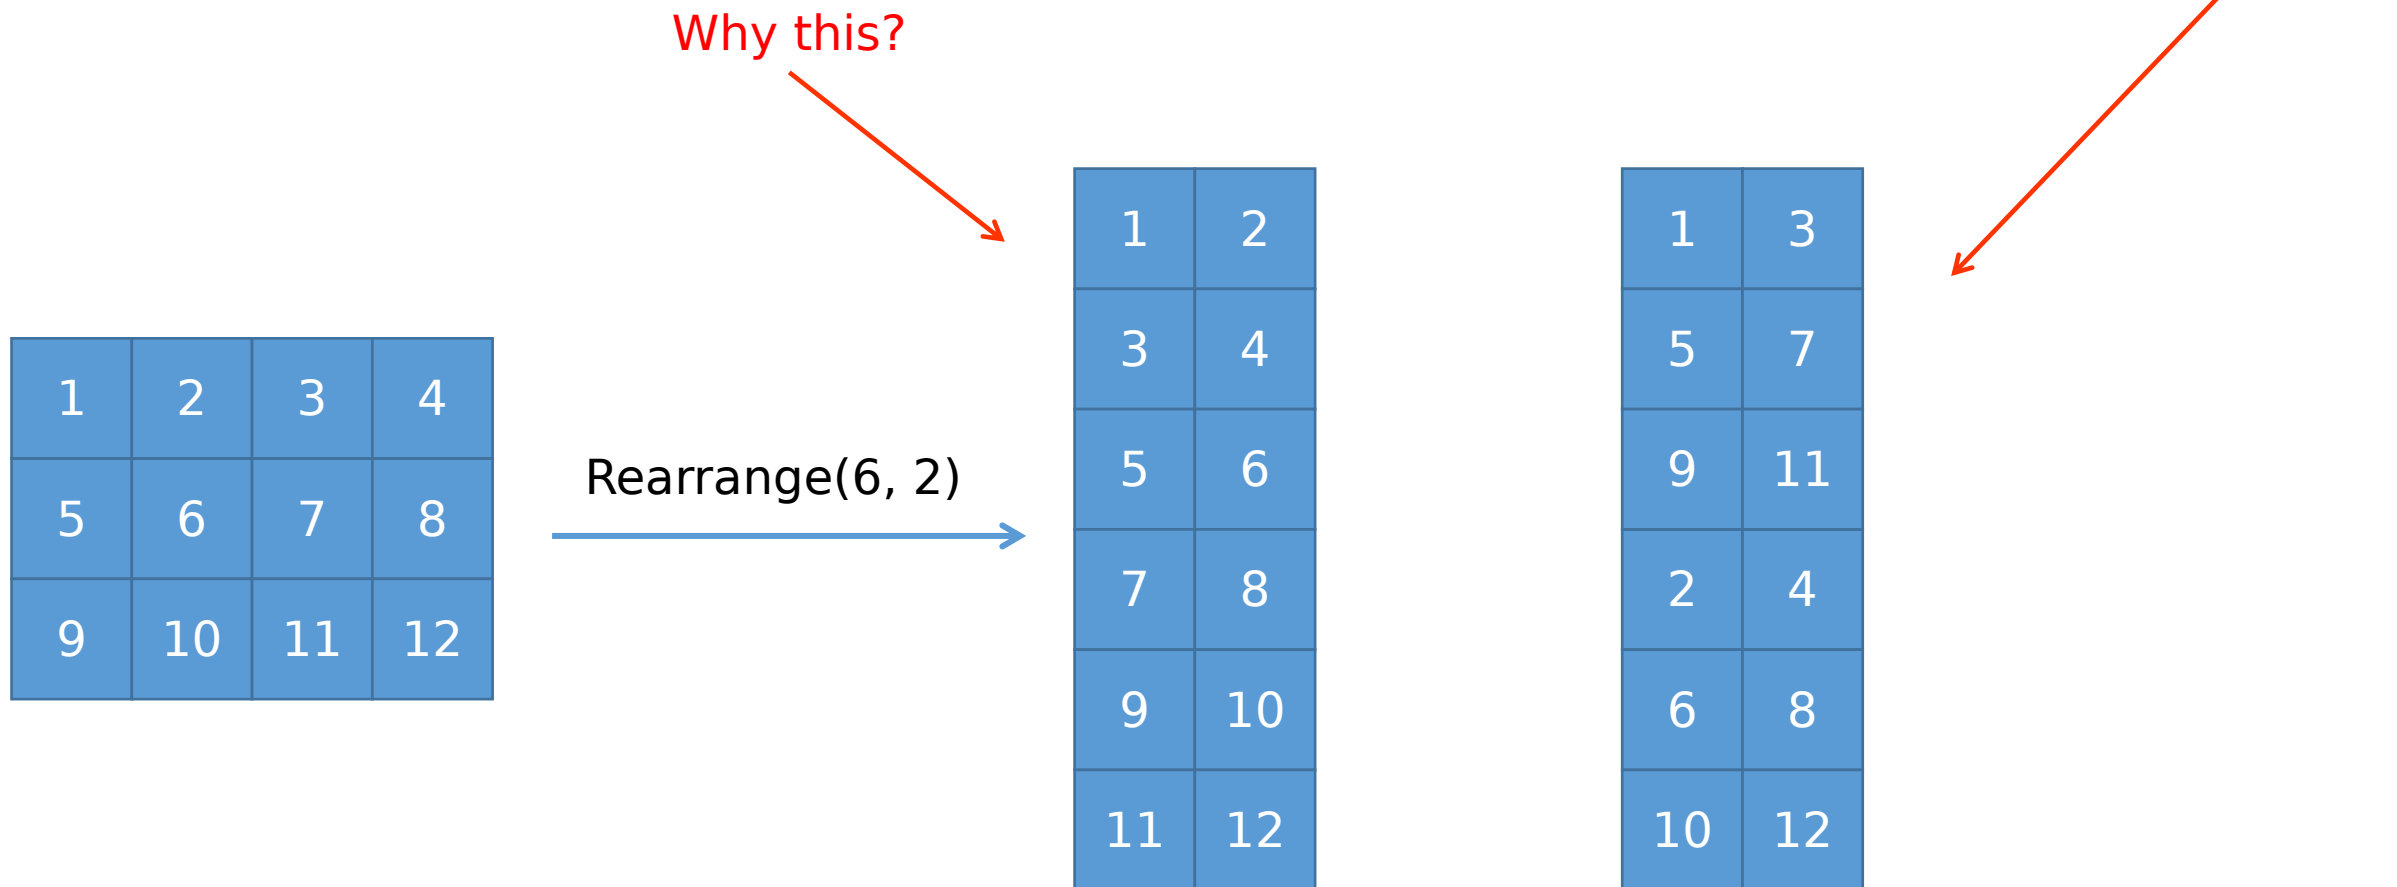


1	2
3	4
5	6
7	8
9	10
11	12

Shape = (6, 2)

Tensor Algebra: Rearrange (2)

Rearranging a tensor **implicitly** defines an **order**.



Tensor Algebra: Rearrange (3)

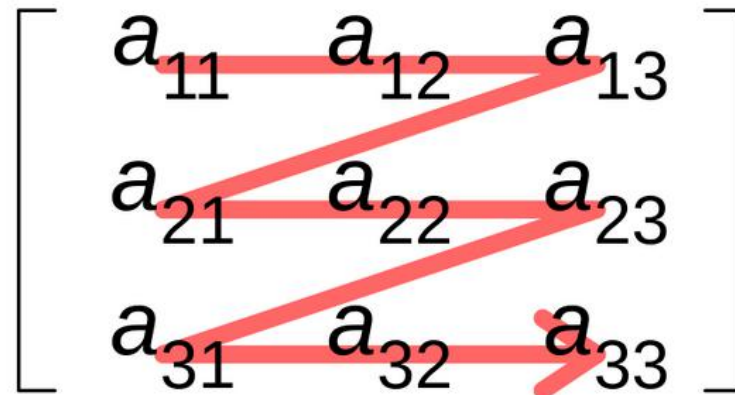
There is **no right answer**.

PyTorch uses **row-major** tensor algebra.

When **rearranging** elements are **ordered** from **left to right starting** from the **first (row) dimension**.

Elements **order** is **maintained** during the **rearrange** operation.

Row-major order



Tensor Algebra: Rearrange (4)

In PyTorch there are **two ways** to **rearrange** a tensor:

- reshape

- **Actually rearrange** the tensor in the **memory**
- Slow rearrange, fast access

- view

- Only **modifies** the **indexing**. The **memory is not modified**
- Fast rearrange, slightly slower access

```
import torch

a = torch.tensor([[1,2,3,4], [5,6,7,8], [9,10,11,12]])
b = a.reshape((6, 2))
c = a.view((6, 2))

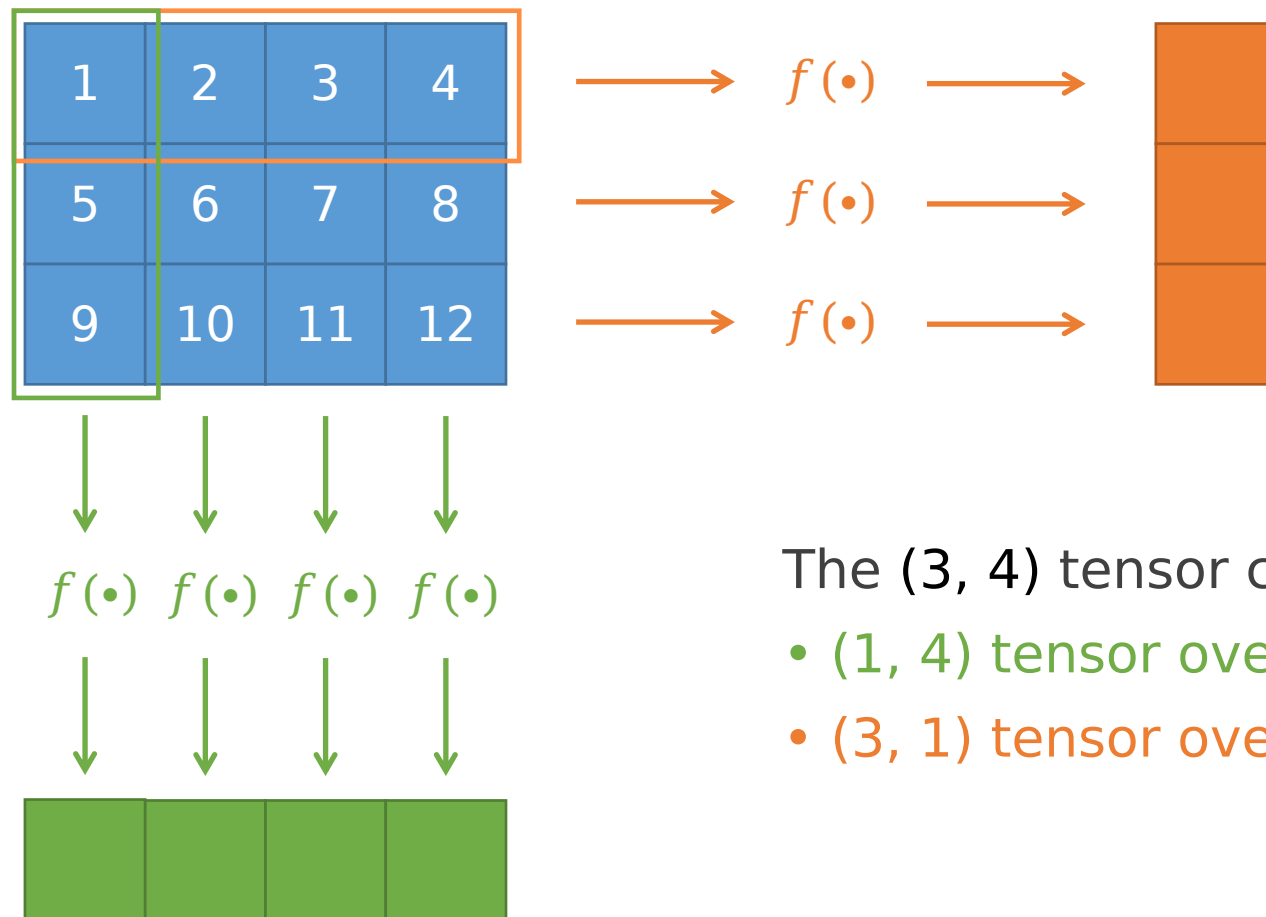
print("b =", b)
print("c =", c)
```



```
↳ b = tensor([[ 1,  2],
              [ 3,  4],
              [ 5,  6],
              [ 7,  8],
              [ 9, 10],
              [11, 12]])
c = tensor([[ 1,  2],
            [ 3,  4],
            [ 5,  6],
            [ 7,  8],
            [ 9, 10],
            [11, 12]])
```

Tensor Algebra: Reduce

In a **reduction** the **elements** of a **dimension** are **combined** together and **reduced** into a **new element**.



The (3, 4) tensor can be reduced to a:

- (1, 4) tensor over the first dimension (rows)
- (3, 1) tensor over the second dimension (columns)

Tensor Algebra: Reduce (2)

There are **several reduction** operations.

For example:

- Max
- Min
- Sum
- Mean
- Std
- ...

dim is the keyword
to specify the reduction
dimension

```
import torch

a = torch.tensor([[1,2,3,4], [5,6,7,8], [9,10,11,12]]).float()
b = a.mean(dim=0)
c = a.mean(dim=1)
d = a.max(dim=0).values

print("b =", b)
print("c =", c)
print("d =", d)
```

```
b = tensor([5., 6., 7., 8.])
c = tensor([ 2.5000,  6.5000, 10.5000])
d = tensor([ 9., 10., 11., 12.] )
```

Tensor Algebra: Matrix Multiplication

Matrix Multiplication is a binary operation that **produces** a **matrix** from **two matrices**.

The number of **columns** in the **first matrix** must be **equal** to the number of **rows** in the **second matrix**.

$$C = A \times B$$

$$C_{ij} = \sum_{k=0}^{n-1} a_{ik} \cdot b_{kj}$$

Tensor Algebra: Matrix Multiplication (2)

```
import torch

a = torch.tensor([[1,2,3], [4,5,6]])
b = torch.tensor([[1,2,3,4], [5,6,7,8], [9,10,11,12]])

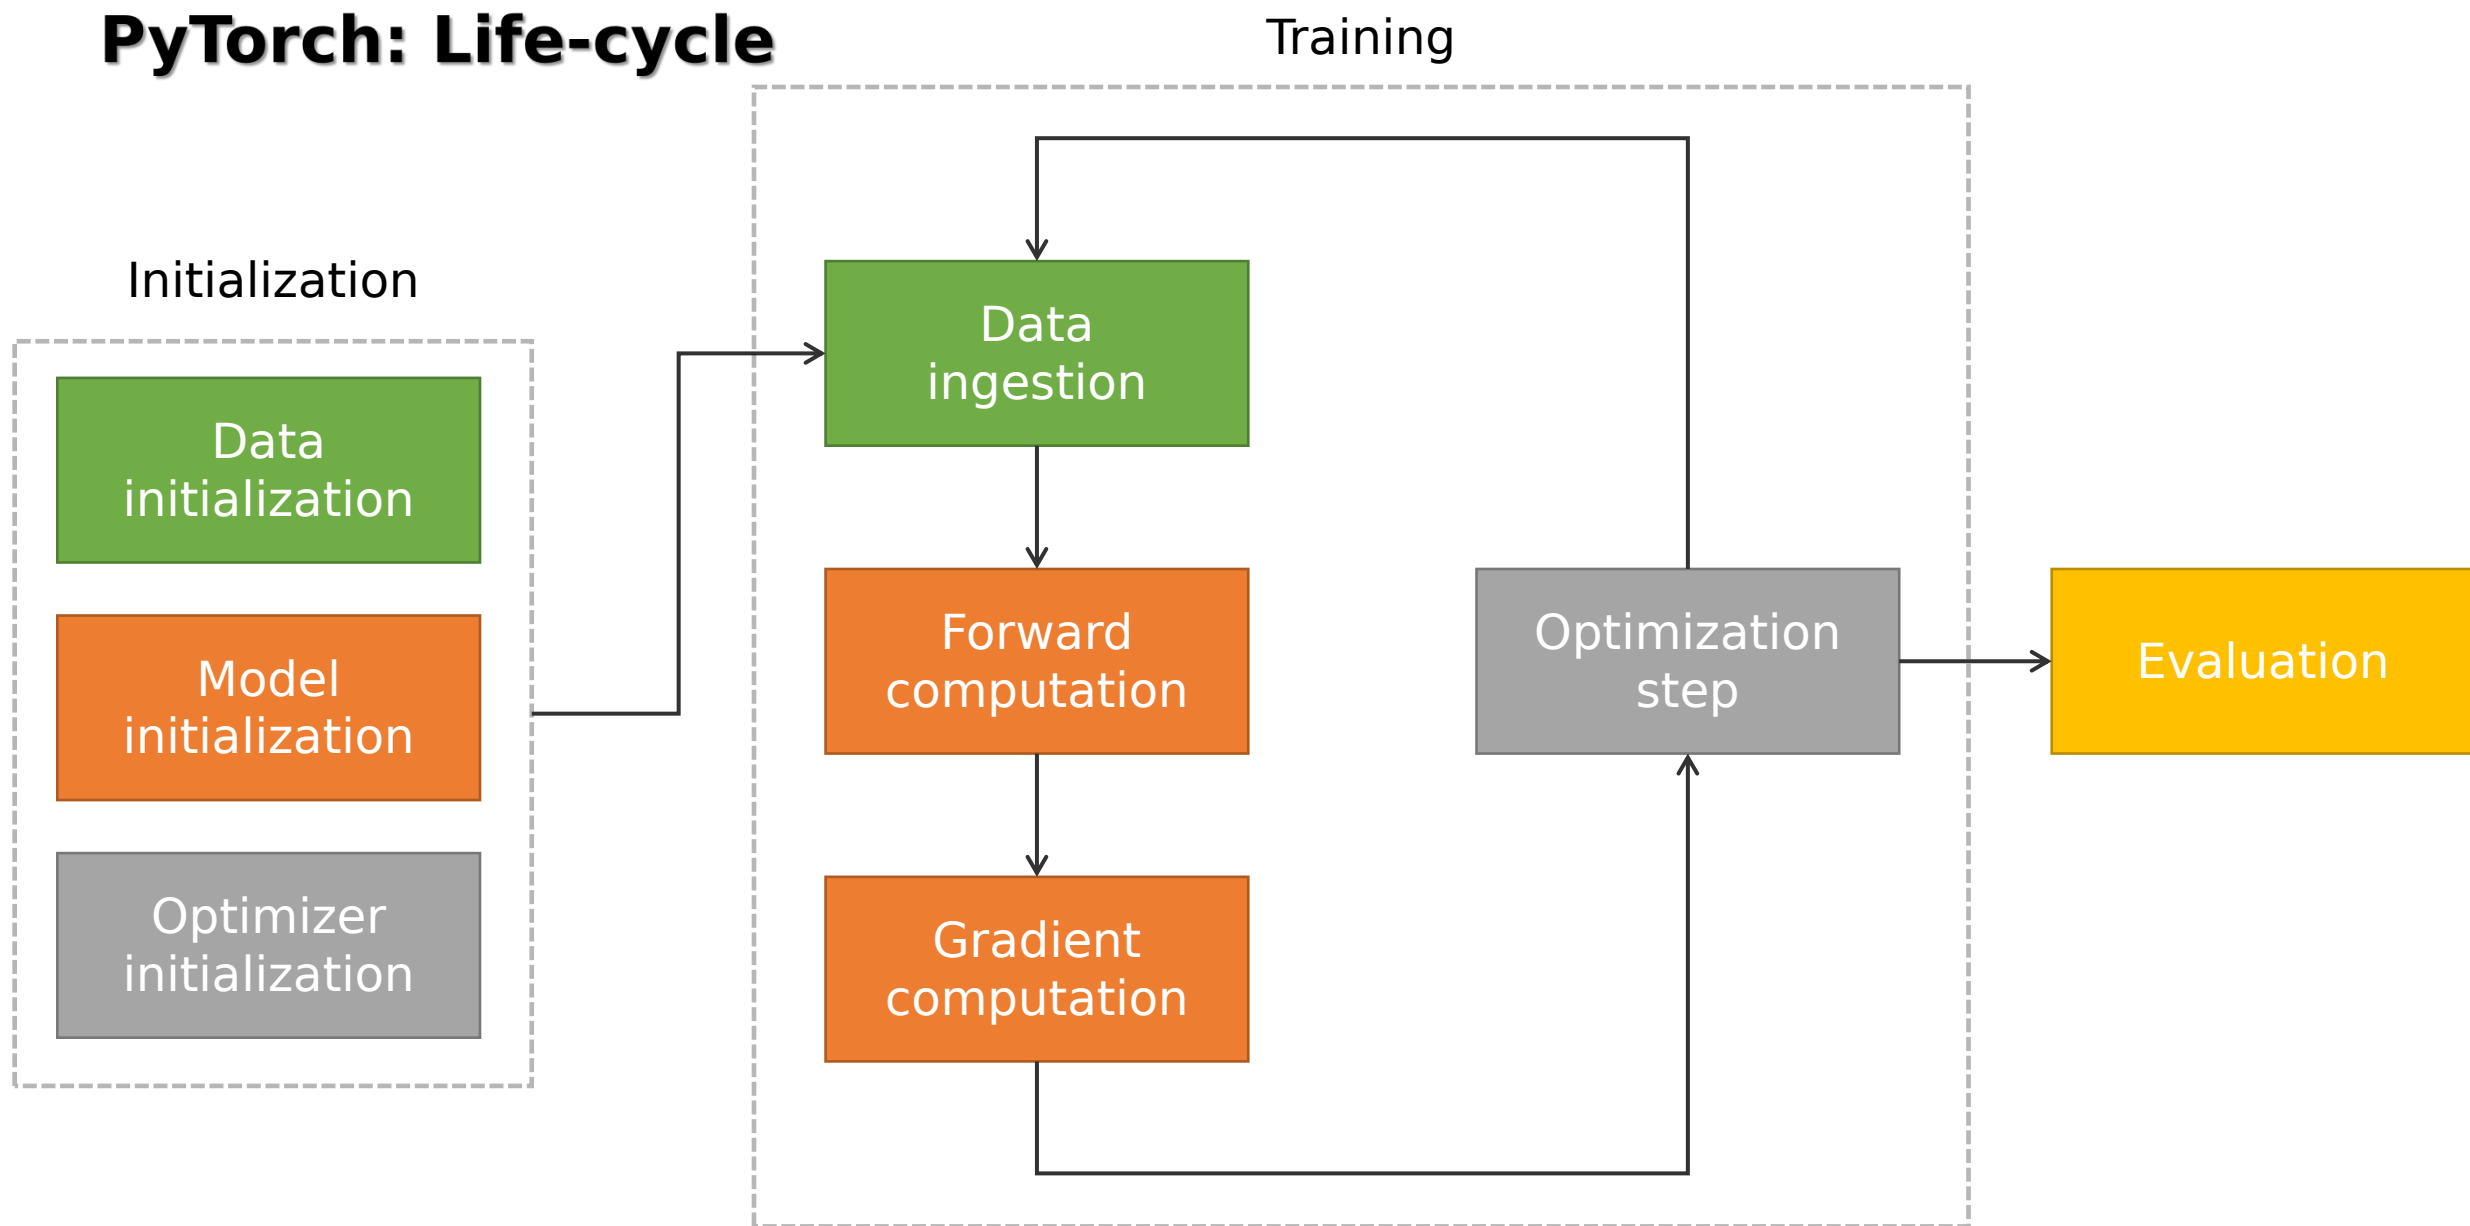
c = torch.matmul(a,b)

print("c = ", c)

print("shape a =", a.shape)
print("shape b =", b.shape)
print("shape c =", c.shape)
```

```
↳ c = tensor([[ 38,  44,  50,  56],
               [ 83,  98, 113, 128]])
shape a = torch.Size([2, 3])
shape b = torch.Size([3, 4])
shape c = torch.Size([2, 4])
```

PyTorch: Life-cycle



PyTorch: Tensors

There are several ways to initialize a tensor in PyTorch:

- From a **list** or a **numpy array**
- Using `torch.ones(shape)` to initialize it with **ones**
- Using `torch.zeros(shape)` to initialize it with **zeros**
- Using the **random module**
- etc..

```
import torch

a = torch.tensor([[1,2,3], [4,5,6]])
b = torch.ones((3,3))
c = torch.rand((2, 3))

print("a =", a)
print("b =", b)
print("c =", c)
```

```
➦ a = tensor([[1, 2, 3],
              [4, 5, 6]])
    b = tensor([[1., 1., 1.],
              [1., 1., 1.],
              [1., 1., 1.]])
    c = tensor([[0.7326, 0.1586, 0.8283],
              [0.3195, 0.4380, 0.5967]])
```

PyTorch: Optimizers

There are **several SGD** optimizers implemented in the **torch.optim** package

For example:

- Vanilla SGD
- Adam
- Adagrad
- RMSProp
- etc...

Algorithms [↗](#)

Adadelta

Adagrad

Adam

AdamW

PyTorch: Functional Interface

The **functional interface** is located in the `torch.nn.functional` **package**.

There are several functions available:

- Activation functions
- Linear functions
- Loss functions
- etc...

`torch.nn.functional`

Convolution functions

Pooling functions

Attention Mechanisms

Non-linear activation functions

Linear functions

Dropout functions

Sparse functions

Distance functions

Loss functions

Vision functions

PyTorch: Dataset

In PyTorch there are two type of dataset abstractions:

- `torch.utils.data.Dataset`
- `torch.utils.data.IterableDataset`

`torch.utils.data.Dataset` follows the **map** python data model.

You have to implement the methods:

- **`__len__(self)`**: returns the number of data points
- **`__getitem__(self, i)`**: returns the **i-th** data point

`torch.utils.data.IterableDataset` follows the **iterable** python data model.

You have to implement the method:

- **`__iter__(self)`**: returns an iterable over the data points

PyTorch: Dataset (2)

```
class SinDataset(torch.utils.data.Dataset):  
    def __init__(self, points):  
        self.points = points  
  
    def __len__(self):  
        return self.points  
  
    def __getitem__(self, idx):  
        norm = (idx / self.points) * (2*torch.pi)  
        return math.sin(norm)
```

```
ds = SinDataset(100)  
print("len:", len(ds))  
print("0:", ds[0])  
print("1:", ds[25])  
print("2", ds[75])
```

```
len: 100  
0: 0.0  
1: 1.0  
2 -1.0
```

PyTorch: Dataset (3)

For **`torch.utils.data.IterableDataset`** it's easier and less error prone to make the dataset class itself an **iterator**. This can be achieved by implementing the methods:

- **`__iter__(self)`** that returns **`self`**
- **`__next__(self)`** that returns the **next data point**

```
class MyIterableDataset(torch.utils.data.IterableDataset):
    def __init__(self, points):
        ...

    def __iter__(self):
        return self

    def __next__(self):
        ...
```


PyTorch: Dataset (4)

```
class UnitCircleDataset(torch.utils.data.IterableDataset):
    def __init__(self, points):
        self.total_points = points
        self.generated_points = 0

    def __iter__(self):
        return self

    def __next__(self):
        if self.generated_points >= self.total_points:
            raise StopIteration

        x = torch.rand(1)
        y = torch.rand(1)
        inside = torch.sqrt(x**2 + y**2) < 1

        self.generated_points += 1
        return x, y, inside

ds = UnitCircleDataset(5)
for elem in ds:
    print(elem)

(tensor([0.6505]), tensor([0.1626]), tensor([True]))
(tensor([0.3634]), tensor([0.9226]), tensor([True]))
(tensor([0.0676]), tensor([0.9742]), tensor([True]))
(tensor([0.6160]), tensor([0.9532]), tensor([False]))
(tensor([0.5621]), tensor([0.5406]), tensor([True]))
```

PyTorch: DataLoader

the class **torch.utils.data.DataLoader**

Combines a dataset and a sampler, and provides an iterable over the given dataset.

The **DataLoader** supports **both map-style** and **iterable-style** datasets with single- or multi-process loading, customizing loading order and automatic **batching**.

```
ds = UnitCircleDataset(100)

dl = torch.utils.data.DataLoader(
    dataset=ds,
    batch_size=4,
)

for batch in dl:
    x, y, inside = batch
    print(type(batch), len(batch))
    print(x.shape)
    print(y.shape)
    print(inside.shape)
    break
```

```
<class 'list'> 3
torch.Size([4, 1])
torch.Size([4, 1])
torch.Size([4, 1])
```

PyTorch: DataLoader (2)

the class **`torch.utils.data.DataLoader`** has various arguments, the most important ones are:

- dataset
 - The dataset which to load the data from
- batch_size
 - The batch_size to use
- shuffle
 - Use a random order when accessing the dataset, works only for map-style datasets
- num_workers
 - How many parallel processes to use, each process loads one data point
- sampler
 - The sampler to use

PyTorch: Module

The fundamental building block of PyTorch models is the **torch.nn.Module** class.

- **Parameters**

- Each model may encompass parameters, represented by tensors wrapped with **torch.nn.Parameter**.
- Any **attribute** of a Module that **is** a **torch.nn.Parameter** is automatically **included** in the **parameters list** of that module.

- **Sub-modules**

- Models can be **composed** of **sub-modules**.
- All **parameters within** these **sub-modules** are seamlessly **aggregated** into the **parameters list** of the **containing module**.

PyTorch: Module (1)

PyTorch **provides** a comprehensive **set of modules** that serve as the **foundational building blocks** for constructing **neural networks**.

Among these, some of the **most frequently utilized** ones include:

- `torch.nn.Linear`:
 - Implements a linear transformation, commonly used for fully connected layers.
- `torch.nn.Sequential`
 - A container module that allows for the orderly arrangement of other modules in a sequential manner.
- `torch.nn.ReLU`
 - Introduces non-linearity through the Rectified Linear Unit (ReLU) activation function.
- `torch.nn.Conv2d` ← next lecture
 - The 2d convolution operation.

PyTorch: Module (2)

The most important methods of a **torch.nn.Module** are:

- **__init__(self)**
 - Constructor method where you define the attributes (parameters and sub-modules) of your module.
 - Remember to call the parent constructor with `super().__init__()`
- **forward(self, input)**
 - Crucial method where the actual computation of the module occurs.
 - You specify the forward pass of your neural network in this function.
 - Receives input data and returns the output of the module.
- **parameters(self)**
 - Returns an iterator over module parameters.
 - Parameters are tensors that are automatically optimized during training.
- **to(self, device)**
 - Moves the module to the specified device (CPU or GPU).
 - Useful for managing device allocation.

PyTorch: Module (3)

```
class QuadraticModel(torch.nn.Module):
    def __init__(self):
        super().__init__()
        self.a = torch.nn.Parameter(torch.rand(1))
        self.b = torch.nn.Parameter(torch.rand(1))
        self.c = torch.nn.Parameter(torch.rand(1))

    def forward(self, x):
        return self.a*x.pow(2) + self.b*x + self.c
```

```
model = QuadraticModel()
print(list(model.parameters()))
print(model(torch.tensor(2)))
```

```
[Parameter containing:
tensor([0.0897], requires_grad=True), Parameter containing:
tensor([0.6999], requires_grad=True), Parameter containing:
tensor([0.1056], requires_grad=True)]
tensor([1.8640], grad_fn=<AddBackward0>)
```

Training a Module

In PyTorch you can use **variants** of the **Stochastic Gradient Descent** algorithm to **train** a module from the **torch.optim** package.

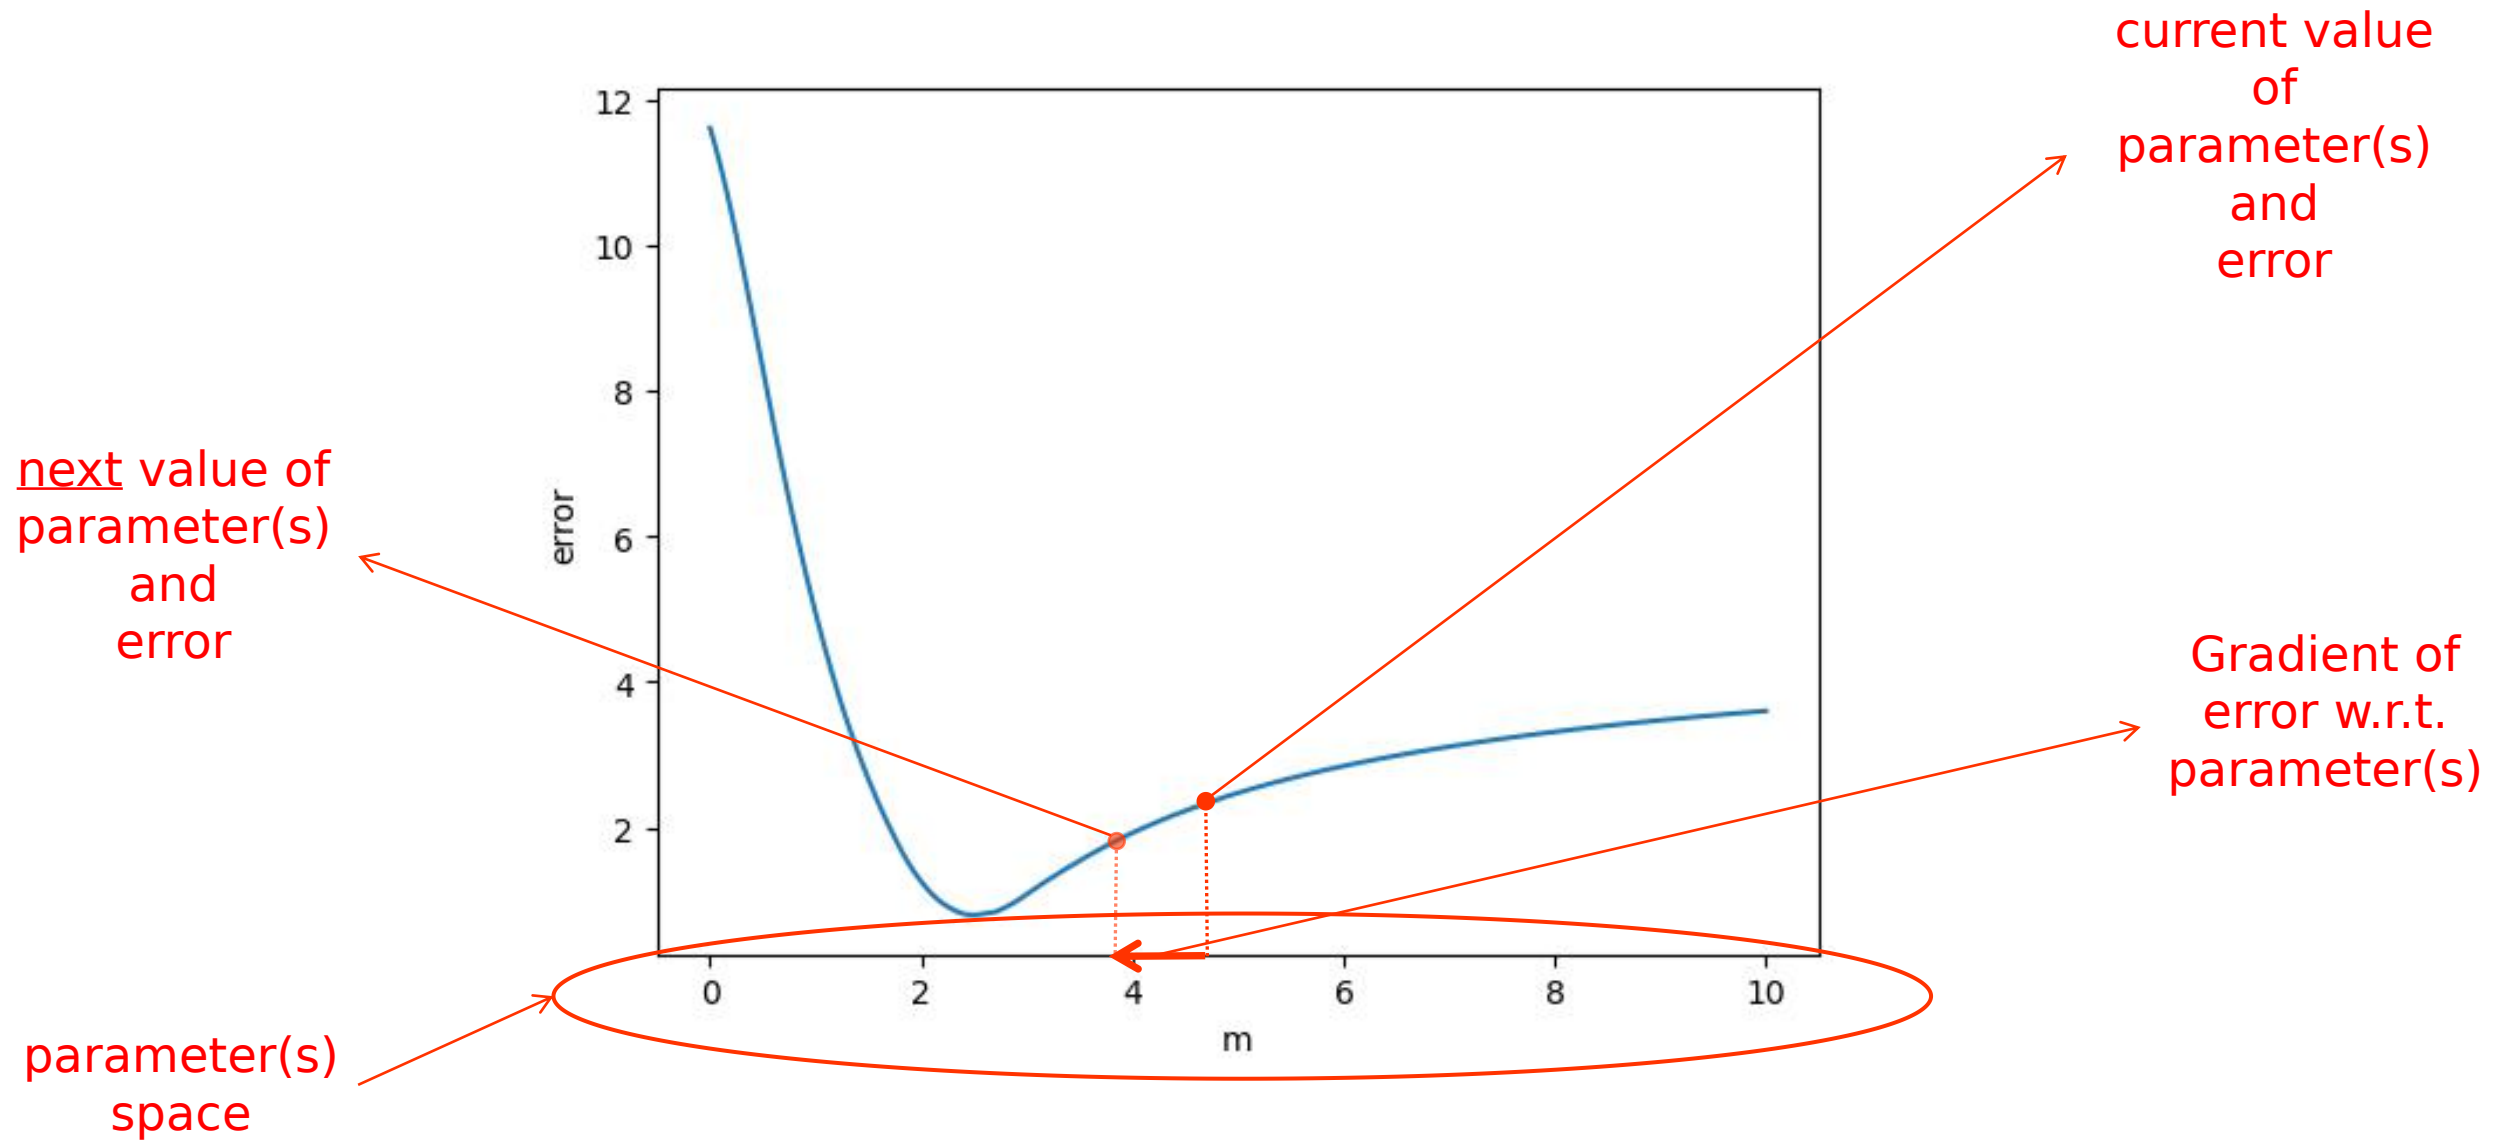
The **main training loop** should (most of the times) **perform** these operations:

- Initialize dataloader, model and optimizer
- for epoch in range(0, epochs):
 - for batch in dataloader
 - output = model(batch.input)
 - error = error_fn(output, batch.target)
 - optimizer.zero_grad()
 - error.backward()
 - optimizer.step()

An “epoch” is training the model over the whole dataset one time.

To train a model you need to perform several epochs

Training a Module (2)



Training a Module (3)

Training a model means **computing an error** (a scalar function), computing the **gradient** of the **error w.r.t.** the model **parameters** and **following** the **direction** that **minimize** the **error**.

PyTorch has some **commonly used** error functions (also called loss functions) already implemented, for example:

- Mean Absolute Error (MAE): **torch.nn.L1Loss**
 - Measures the average absolute difference between predicted and actual values.
 - Less sensitive to outliers compared to MSE.
- Mean Squared Error (MSE): **torch.nn.MSELoss**
 - Measures the average squared difference between predicted and actual values.
 - Suitable for tasks like predicting numerical values.
- Cross-Entropy: **torch.nn.CrossEntropyLoss**
 - Ideal for problems where each sample belongs to one class.
 - Measures the cross-entropy between predicted and target distributions

Exercise 1: (1)

Write the **PointsDataset** class that implements the **torch.utils.data.Dataset** interface:

- Reads a txt file in which each line represents a bidimensional data point, with each dimension separated by a space.
- Saves the content of the file in a data structure of your choice
- The **__len__(self)** method should return the number of data points
- The **__getitem__(self, i)** method should return the i-th data point as a tuple.

```
ds = PointsDataset("dataset.txt")  
print(ds[0])
```

✓ 0.0s

```
(-2.067504630593728, -4.3940160659490966)
```

≡ dataset.txt

```
1 -2.067504630593728 -4.3940160659490966  
2 -9.220342264640013 -26.542222567962877  
3 -3.354104007299002 -9.130317428586983  
4 2.190182077823536 6.24970978934776  
5 8.800144432229736 25.615330535203725  
6 1.2955720622286666 4.457587948472233  
7 -7.368560590563282 -24.042007327266063  
8 -7.885062936720201 -21.902512966035843  
9 0.247269717481778 22.232946480087413
```

Exercise 1: (2)

Write the **LineModule** class that implements the `torch.nn.Module` interface implementing the function $f(x) = wx$.

- LineModule has 1 parameter `w`
- The **`forward(self, x)`** method should return wx

```
model = LineModule()
print(list(model.parameters()))
print(model(torch.tensor([1.])))
```

✓ 0.0s

```
[Parameter containing:
tensor([-0.0583], requires_grad=True)]
tensor([-0.0583], grad_fn=<MulBackward0>)
```

Exercise 1: (3)

Write a complete python script that trains **LineModule** to approximate the data in **dataset1.txt**

- Use the **SGD** optimizer from **torch.optim**
- Use the **MSELoss** from **torch.nn**
- Use a batch_size of 8
- Use a learning rate of 0.001
- Train for 1000 epochs

```
Epoch 0: loss 167.87667012832455
Epoch 0: loss 101.41640371214211
Epoch 0: loss 173.47983306483772
Epoch 0: loss 124.895664870617
Epoch 0: loss 99.13100804449708
Epoch 0: loss 147.6078938833903
Epoch 0: loss 108.22731560725815
Epoch 0: loss 75.81846755449054
Epoch 0: loss 96.93804414637638
Epoch 0: loss 77.51460574750205
Epoch 0: loss 20.281682954088314
Epoch 0: loss 133.66231409443913
Epoch 0: loss 62.41936717002218
Epoch 1: loss 38.00877930294306
Epoch 1: loss 37.03041242639919
```

Exercise 2 (hard)

Train a polynomial model in this form:

$$ax^4 + bx^3 + cx^2 + dx + e$$

To divide the points of **dataset2.txt** (in the form “x y class” each line)

```
dataset2.txt
1 0.6694837762958694 0.8945464835833724 0
2 -0.025934533700649937 1.1107433076159958 0
3 0.9953308180872537 0.25066986949731196 0
4 -0.9322951409309891 0.5909418115823066 0
5 0.16912470429691373 1.0438085765110712 0
6 0.6182140008161767 -0.6014423583073267 1
7 0.5906876196569684 -0.31555239872653296 1
8 -0.6571244845423326 0.7645023530709449 0
9 0.5672735172344342 -0.36143767441115565 1
```


Exercise 2 (hard)

- The polynomial model takes x as input and gives \bar{y} as output
- Given an (x, y) pair from the dataset you can compute $\bar{y} = \text{model}(x)$
- If y is **above** the line (so $y > \bar{y}$) and **it should be** (class 0) then everything is **ok**
- if y is **above** the line and **it should be below it** (class 1) then you should compute an **error**
- And so on...

