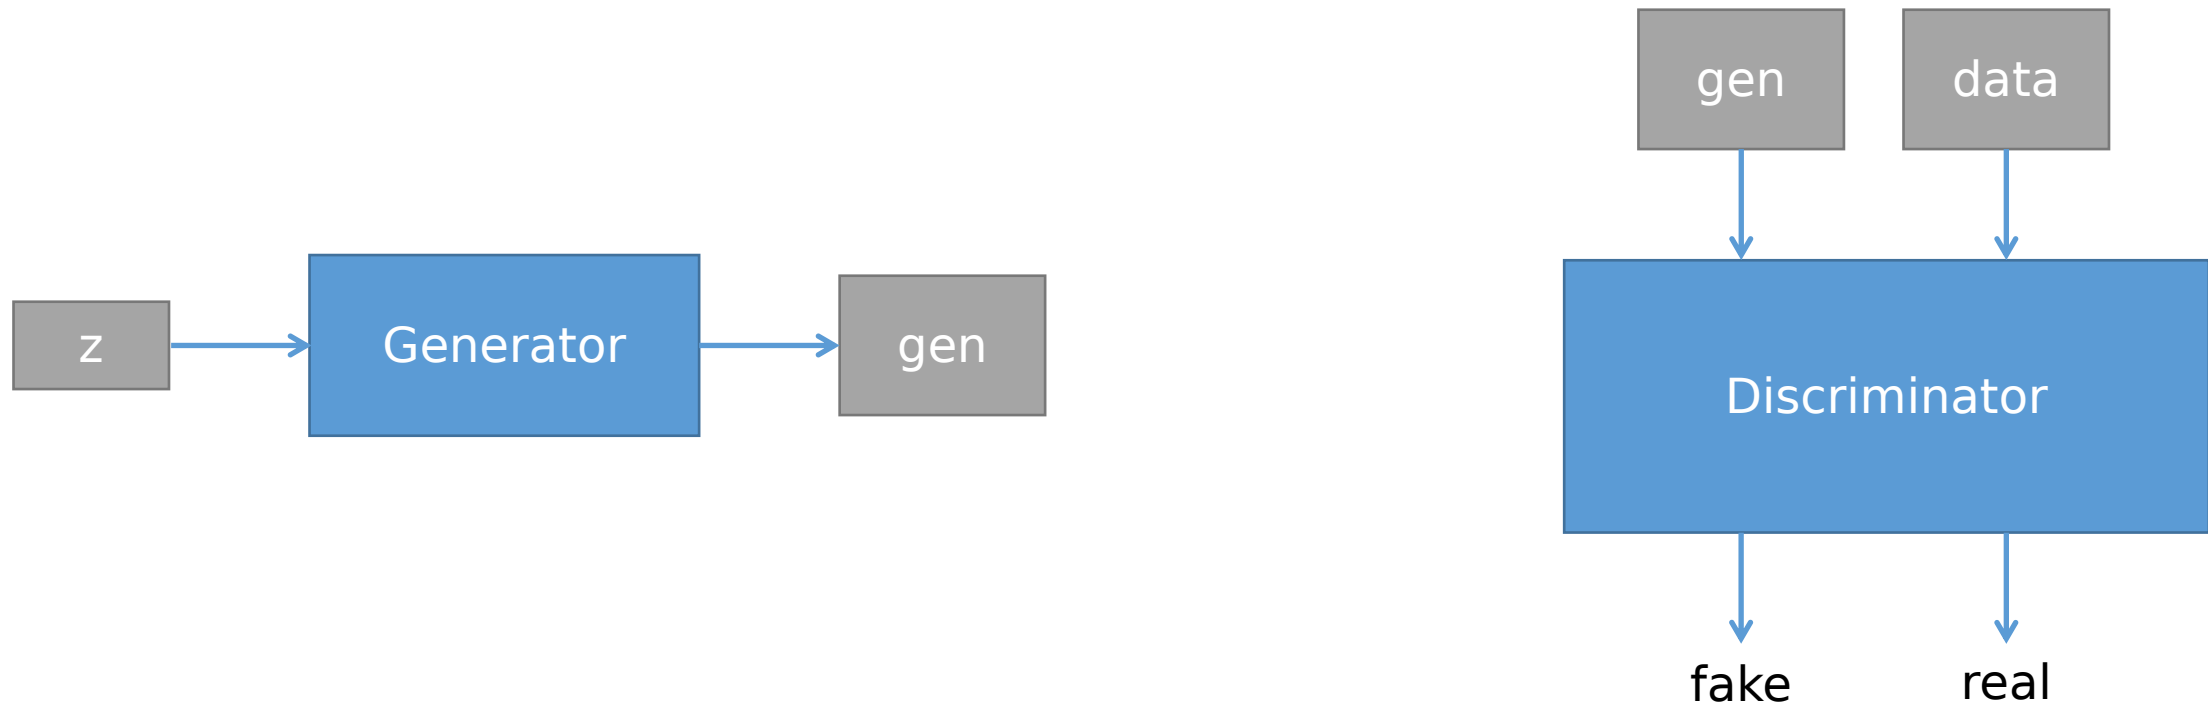


GAN

Generative Adversarial Networks (GANs)

Generative Adversarial Networks (GANs) are a **type** of deep learning **model** composed of **two** neural **networks**: a **generator** and a **discriminator**.

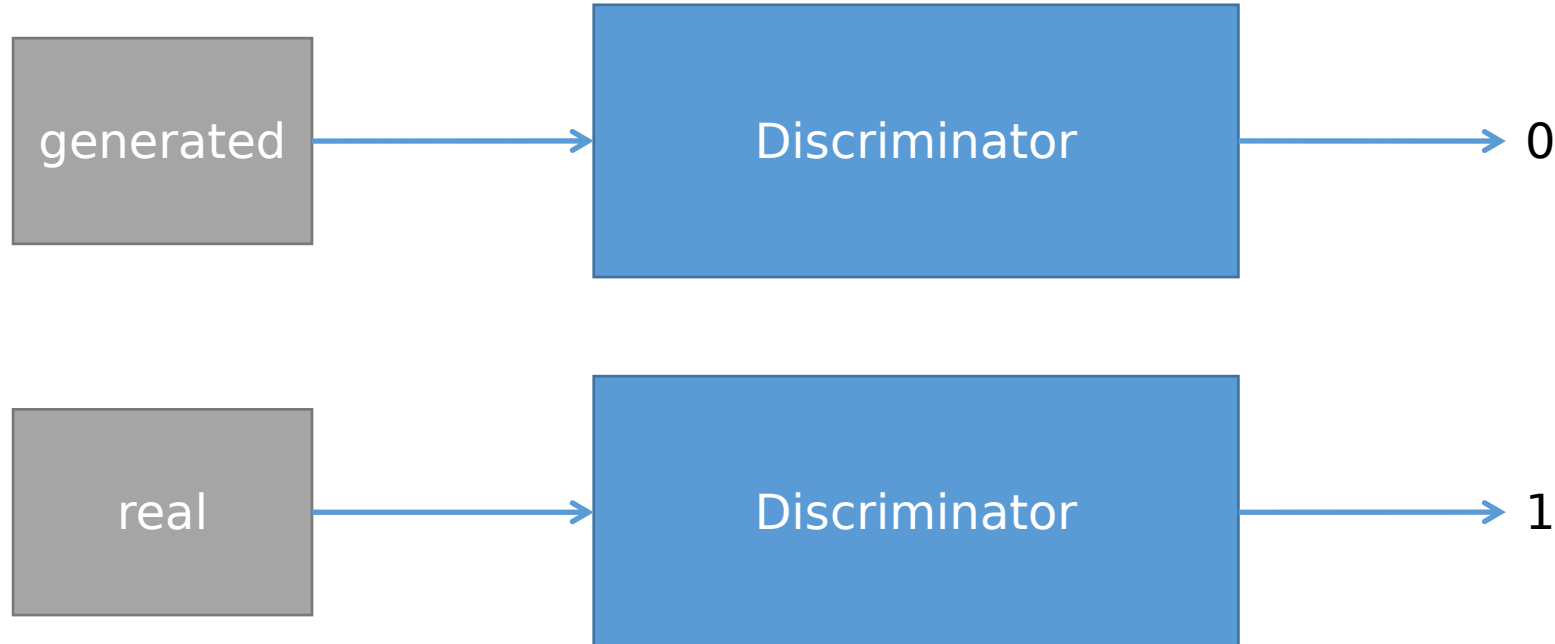
GANs are **designed** for **generating** data that closely **resembles** the **one** in the **training set**.



GANs: Discriminator

The **discriminator** is a **binary classifier** with **one** output.


The **discriminator** is **trained** to **distinguish** between **real** images (output 1) and **generated** images (output 0)



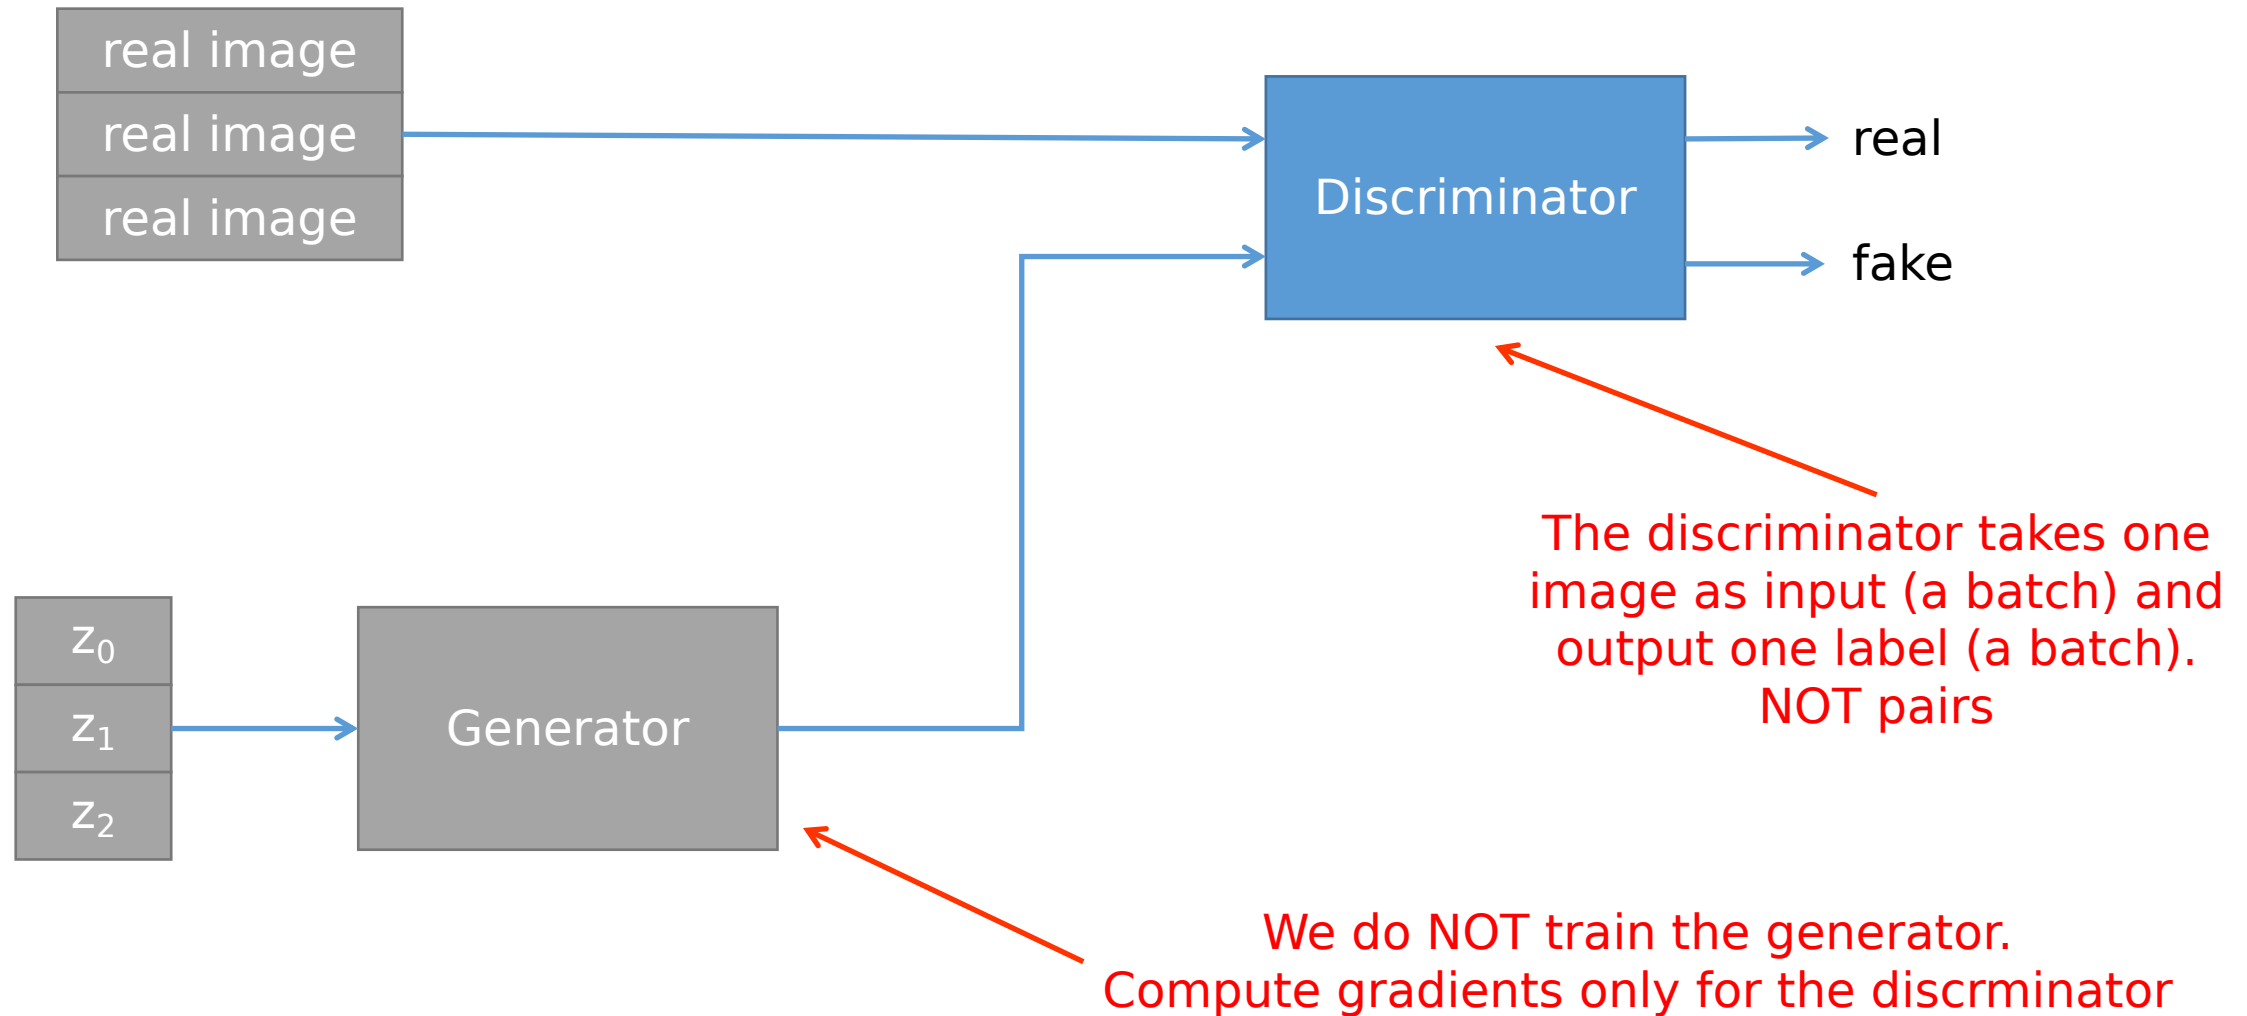
GANs: Generator

The **generator** is trained to **produce** images that can **deceive** the **discriminator**.

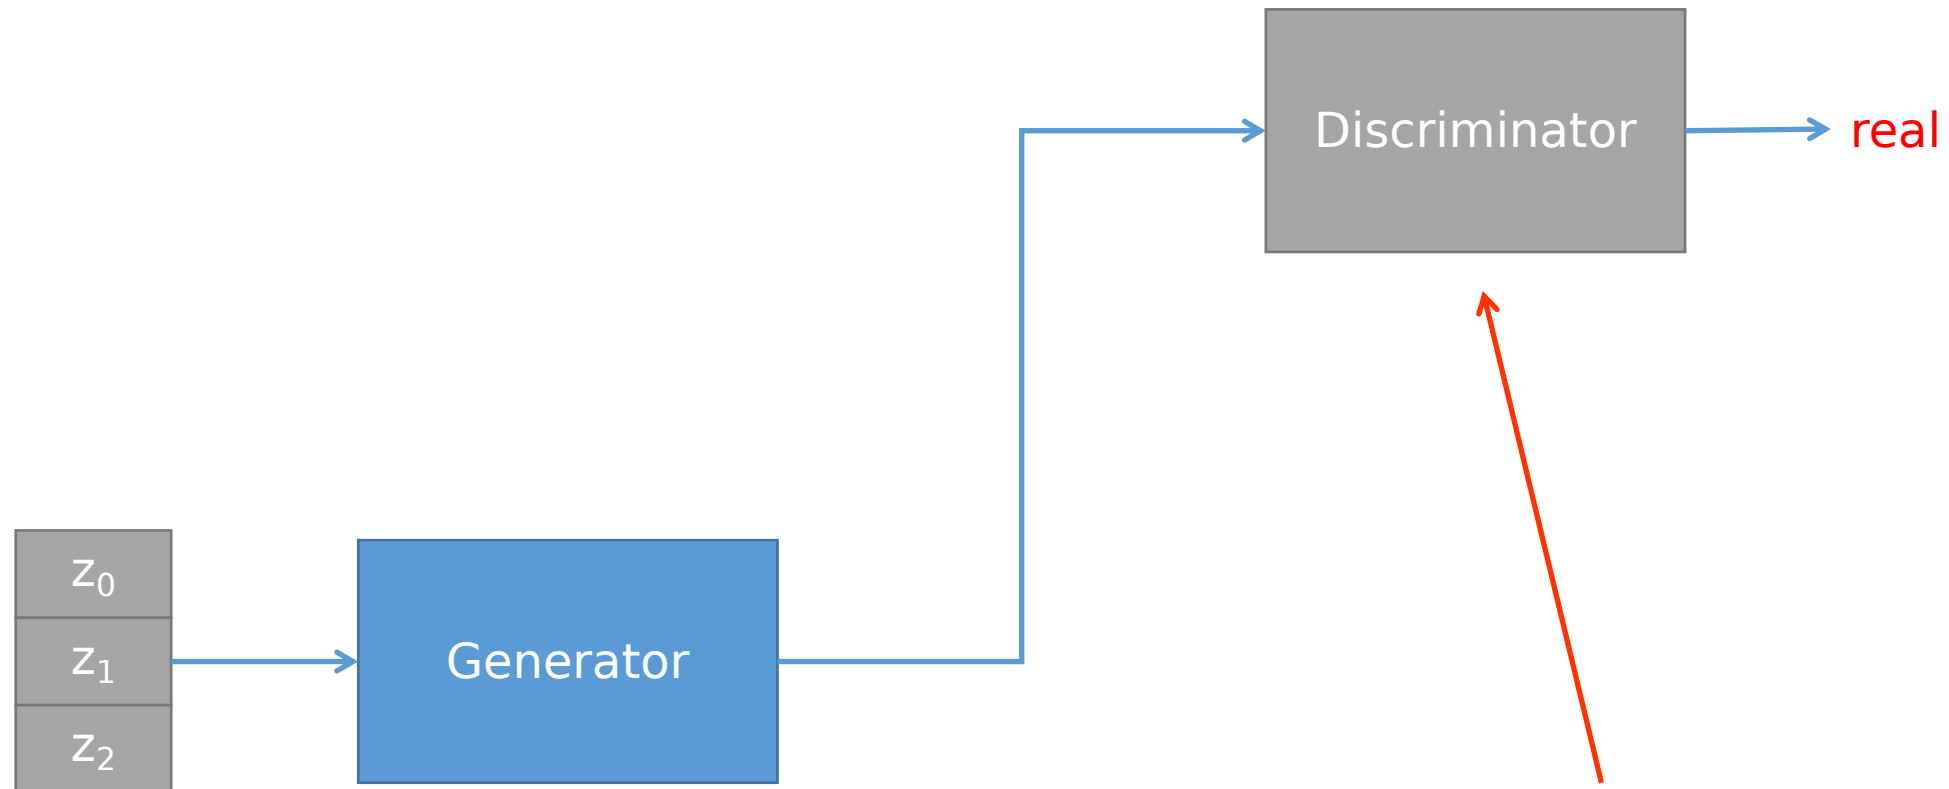
It takes **random** noise as **input** and aims to **generate** an **output** that, when **presented** to the **discriminator**, is **classified** as 'real'.



GANs: Discriminator Training

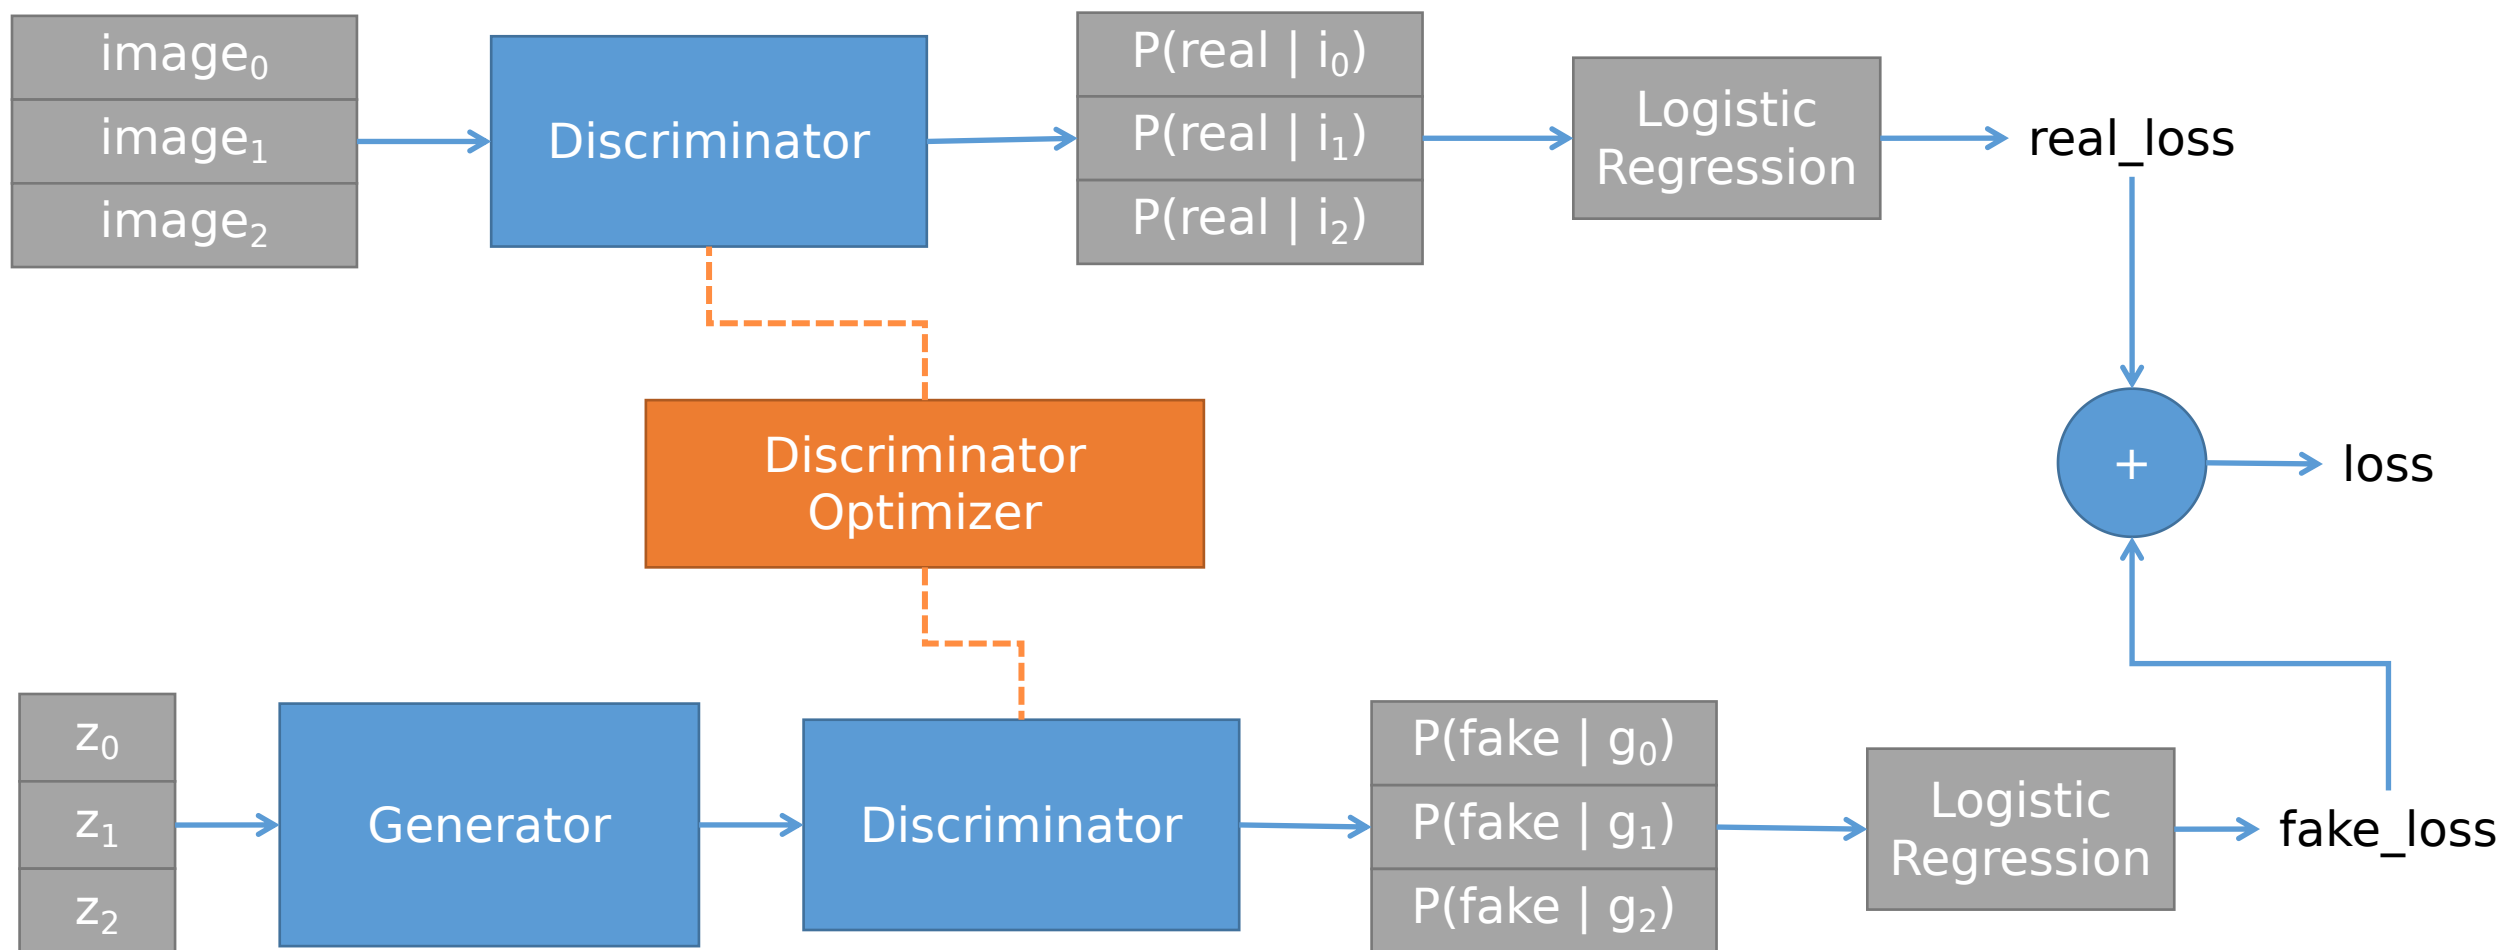


GANs: Generator Training

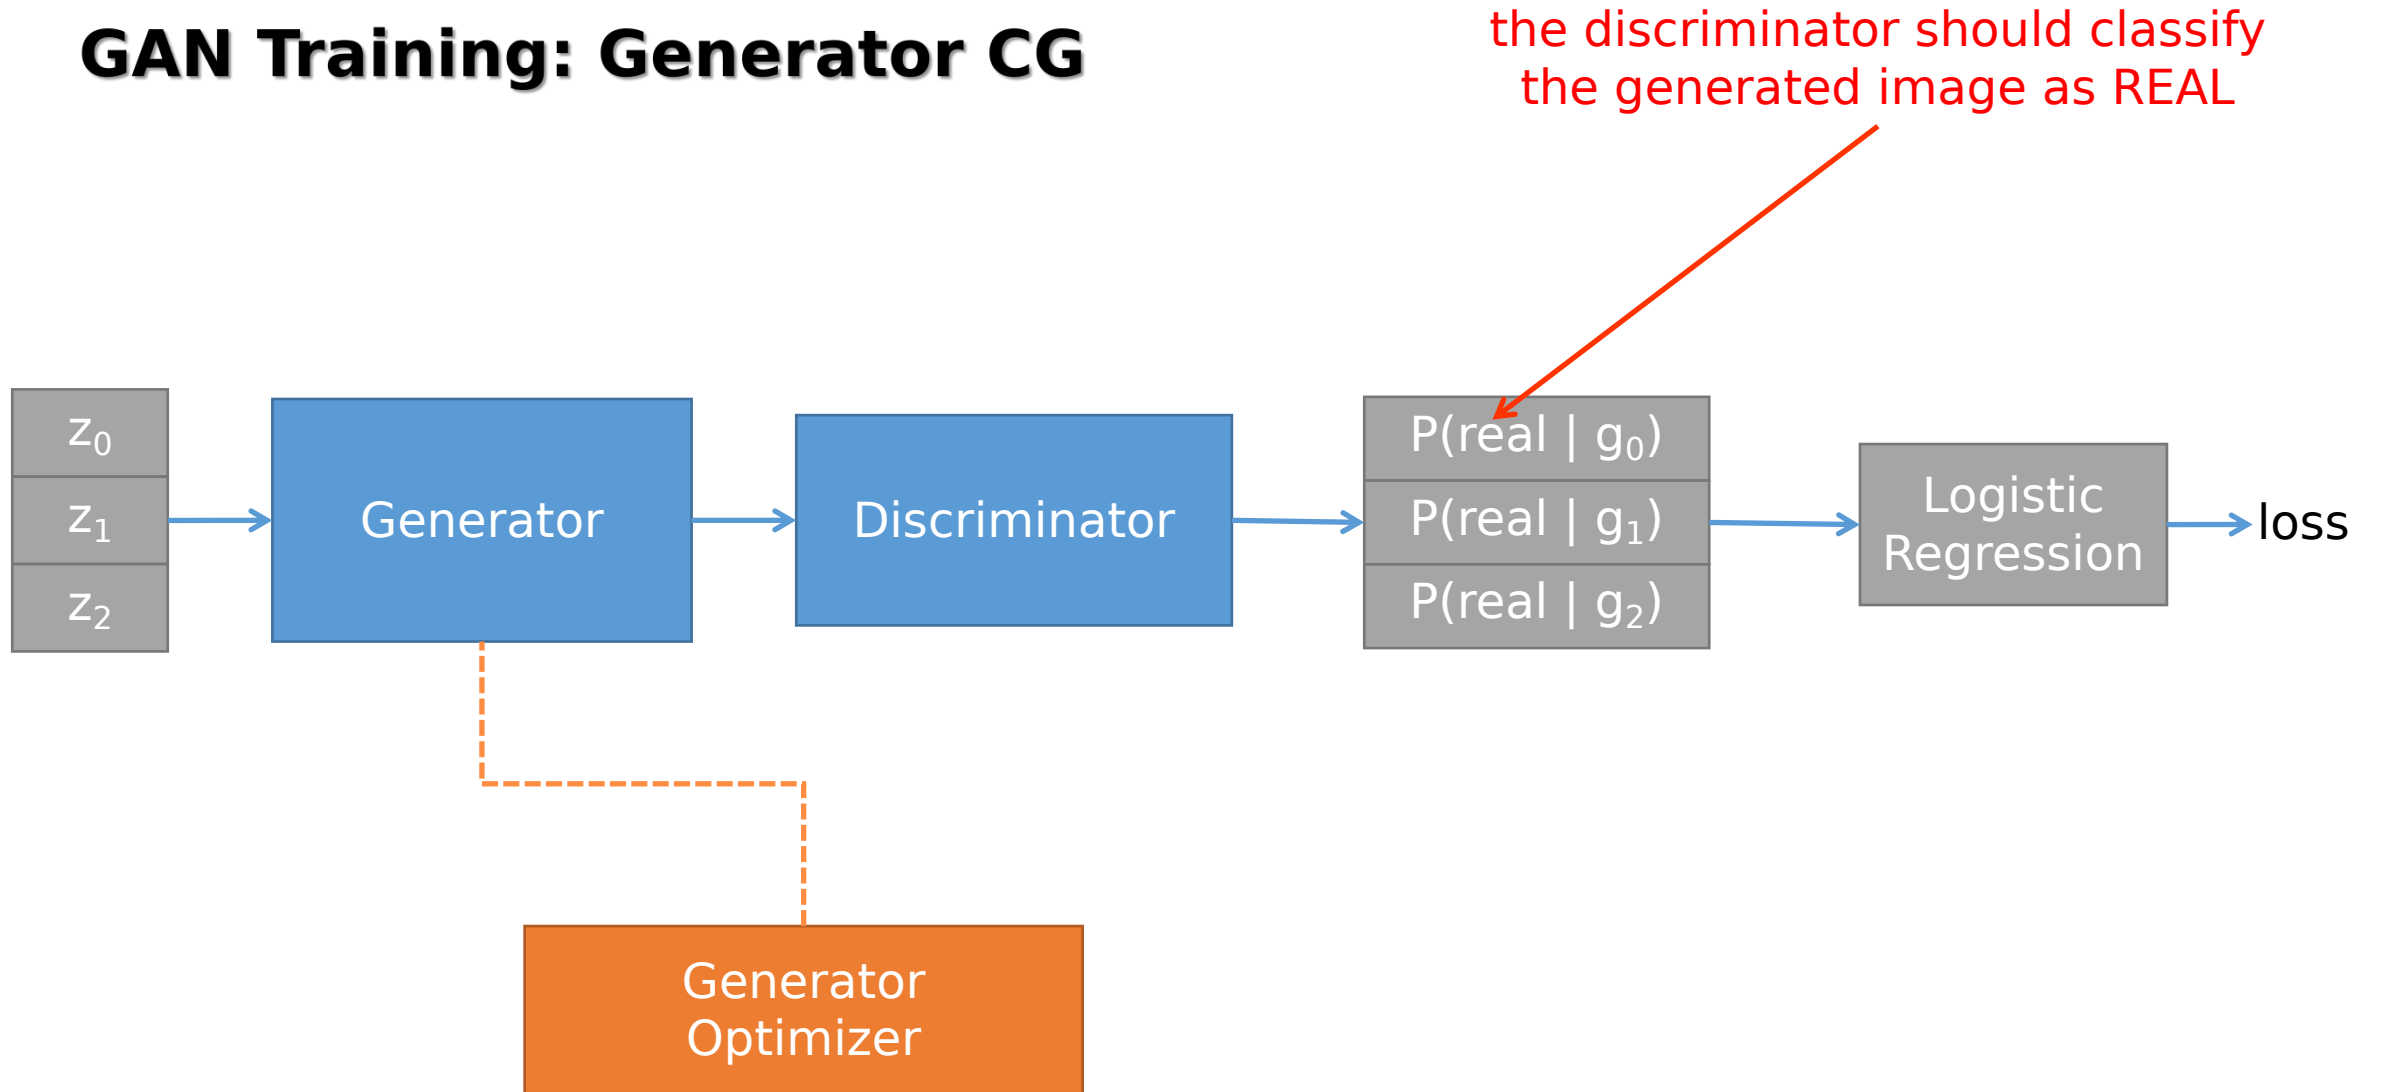


We do NOT train the discriminator.
Compute gradients only for the generator

GAN Training: Discriminator CG



GAN Training: Generator CG



GANs: Discriminator training in practice

To **train** a **GAN** network we use **two optimizer: one for each network**.

To train the **discriminator**:

- randomly **sample** some **real data** from the **dataset**
- **generate** some **fake data** using the **generator**
- **feed** the **discriminator** with the **real** data and **compute** the **error** of the **predicted** label w.r.t. **1** (1 is the real label)
- **feed** the **discriminator** with the **fake** data and **compute** the **error** of the **predicted** labels w.r.t. **0** (0 is the fake label)
- **compute** the **total error** as the **mean** of the two **errors**
- **compute** the **gradients** of the **error** w.r.t. the **discriminator weights**
- **update** the **discriminator weights**

GANs: Generator training in practice

To train the **generator**:

- **generate** some **fake data** using the **generator**
- **feed** the **discriminator** with the **fake** data and **compute** the **error** of the **predicted** labels w.r.t. **1** (1 is the real label)
- **compute** the **gradients** of the **error** w.r.t. the **generator weights**
- **update** the **generator weights**

GANs: PyTorch (initialization)

```
generator = Generator(latent_dim=64, output_dim=28)
discriminator = Discriminator(input_dim=28)
```

```
generator_optimizer = torch.optim.Adam(generator.parameters(), lr=0.0002)
discriminator_optimizer = torch.optim.Adam(discriminator.parameters(), lr=0.0002)
```

GANs: PyTorch (training loop)

```
for epoch in range(0, 20):
    for real_images, _ in tqdm(train_dl):
        real_images = real_images.to("cuda")

        discriminator_optimizer.zero_grad()
        generator_optimizer.zero_grad()
        train_discriminator(
            generator=generator,
            discriminator=discriminator,
            optimizer=discriminator_optimizer,
            real_images=real_images
        )

        discriminator_optimizer.zero_grad()
        generator_optimizer.zero_grad()
        train_generator(
            generator=generator,
            discriminator=discriminator,
            optimizer=generator_optimizer,
            real_images=real_images
        )
```

GANs: PyTorch (discriminator training)

```
def train_discriminator(*, optimizer, generator, discriminator, real_images):  
    batch_size = real_images.shape[0]  
    device = real_images.device  
  
    z = torch.randn(batch_size, generator.latent_dim).to(device)  
    fake_images = generator(z)  
  
    fake_label = torch.zeros(batch_size).to(device)  
    real_label = torch.ones(batch_size).to(device)  
  
    fake_logits = discriminator(fake_images)[: , 0]  
    real_logits = discriminator(real_images)[: , 0]  
  
    fake_loss = torch.nn.functional.binary_cross_entropy(fake_logits, fake_label)  
    real_loss = torch.nn.functional.binary_cross_entropy(real_logits, real_label)  
  
    loss = (fake_loss + real_loss) / 2  
  
    loss.backward()  
    optimizer.step()  
  
    fake_accuracy = (fake_logits < 0.5).float().mean()  
    real_accuracy = (real_logits >= 0.5).float().mean()  
  
    return fake_loss.item(), real_loss.item(), fake_accuracy.item(), real_accuracy.item()
```

GANs: PyTorch (generator training)

```
def train_generator(*, optimizer, generator, discriminator, real_images):  
    batch_size = real_images.shape[0]  
    device = real_images.device  
  
    z = torch.randn(batch_size, generator.latent_dim).to(device)  
    fake_images = generator(z)  
  
    real_label = torch.ones(batch_size).to(device)  
  
    fake_logits = discriminator(fake_images)[:, 0]  
  
    loss = torch.nn.functional.binary_cross_entropy(fake_logits, real_label)  
  
    loss.backward()  
    optimizer.step()  
  
    accuracy = (fake_logits >= 0.5).float().mean()  
  
    return loss.item(), accuracy
```


Auxiliary Classifier GAN

Auxiliary Classifier GAN (ACGAN) is a **variant** of the standard **GAN architecture** that **enhances** generative capabilities by **integrating** an **auxiliary classifier** into the **discriminator**.



It is **ideal** for **conditional image generation** and other tasks requiring **class-specific** data synthesis.

ACGANs: Generator Training


random classes



random noise




random
classes





real

ACGANs: Discriminator Training (generated images)

random classes

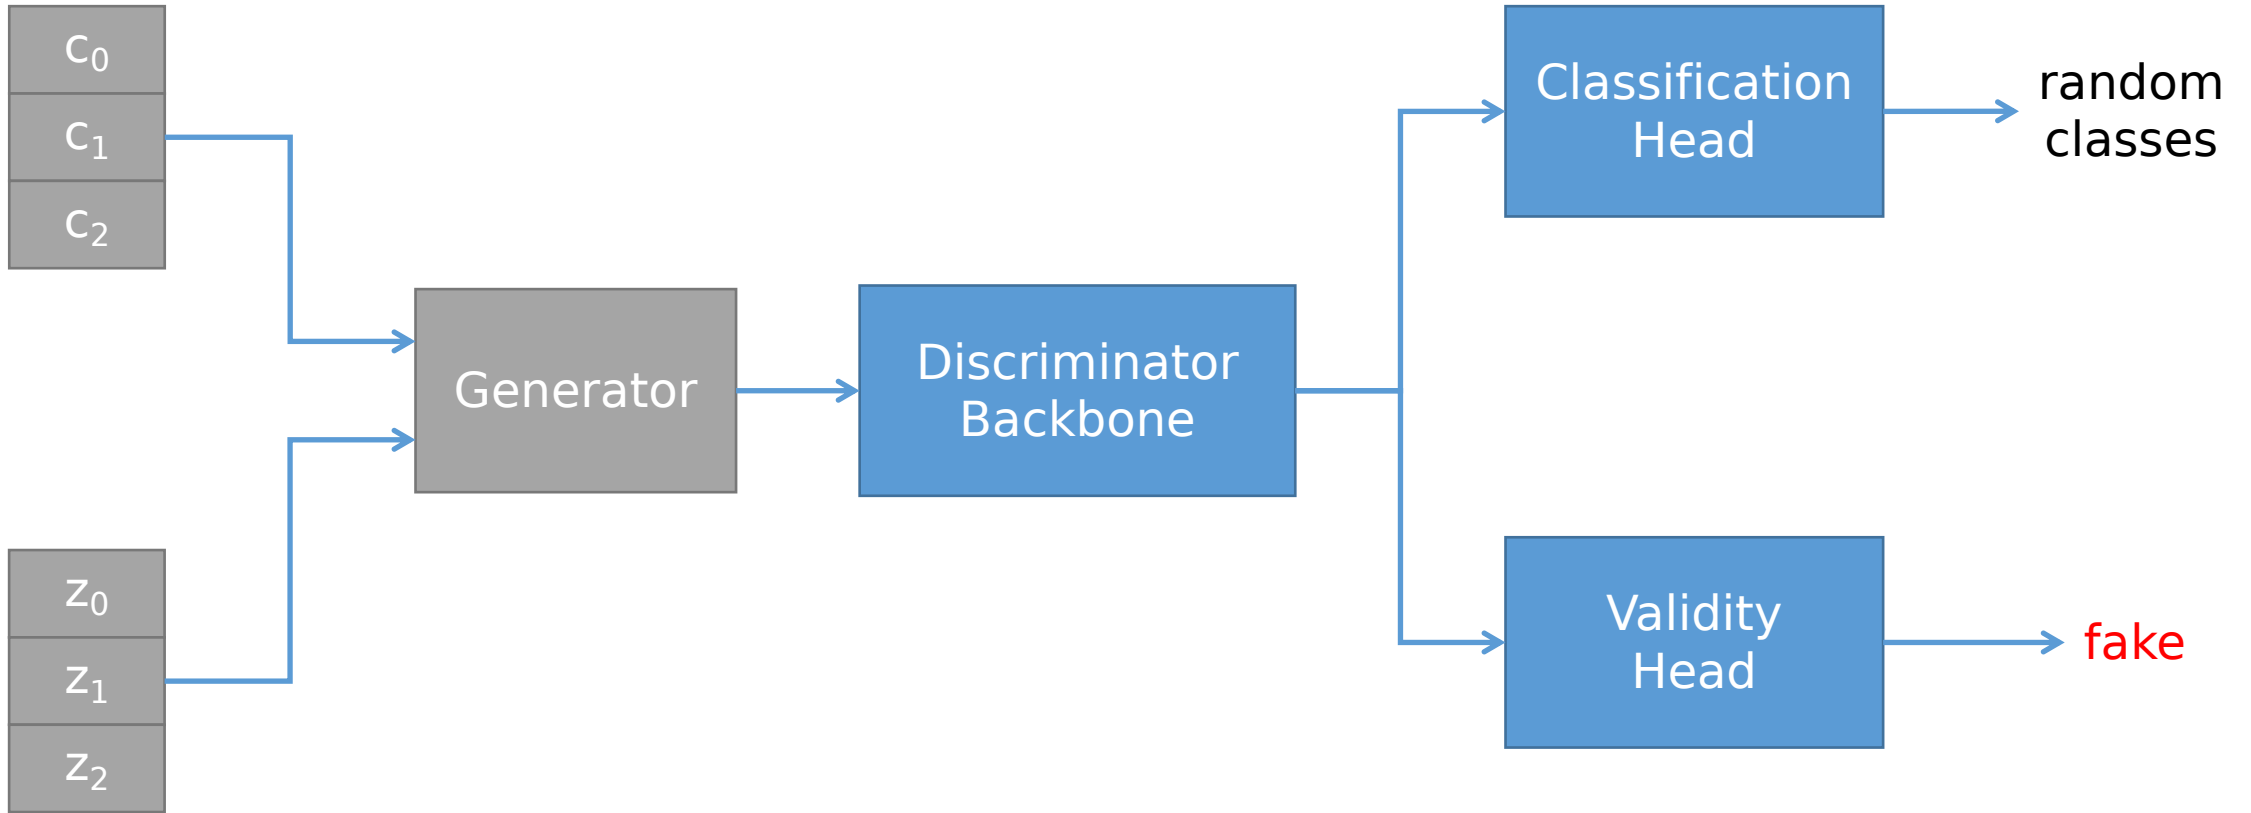


random noise

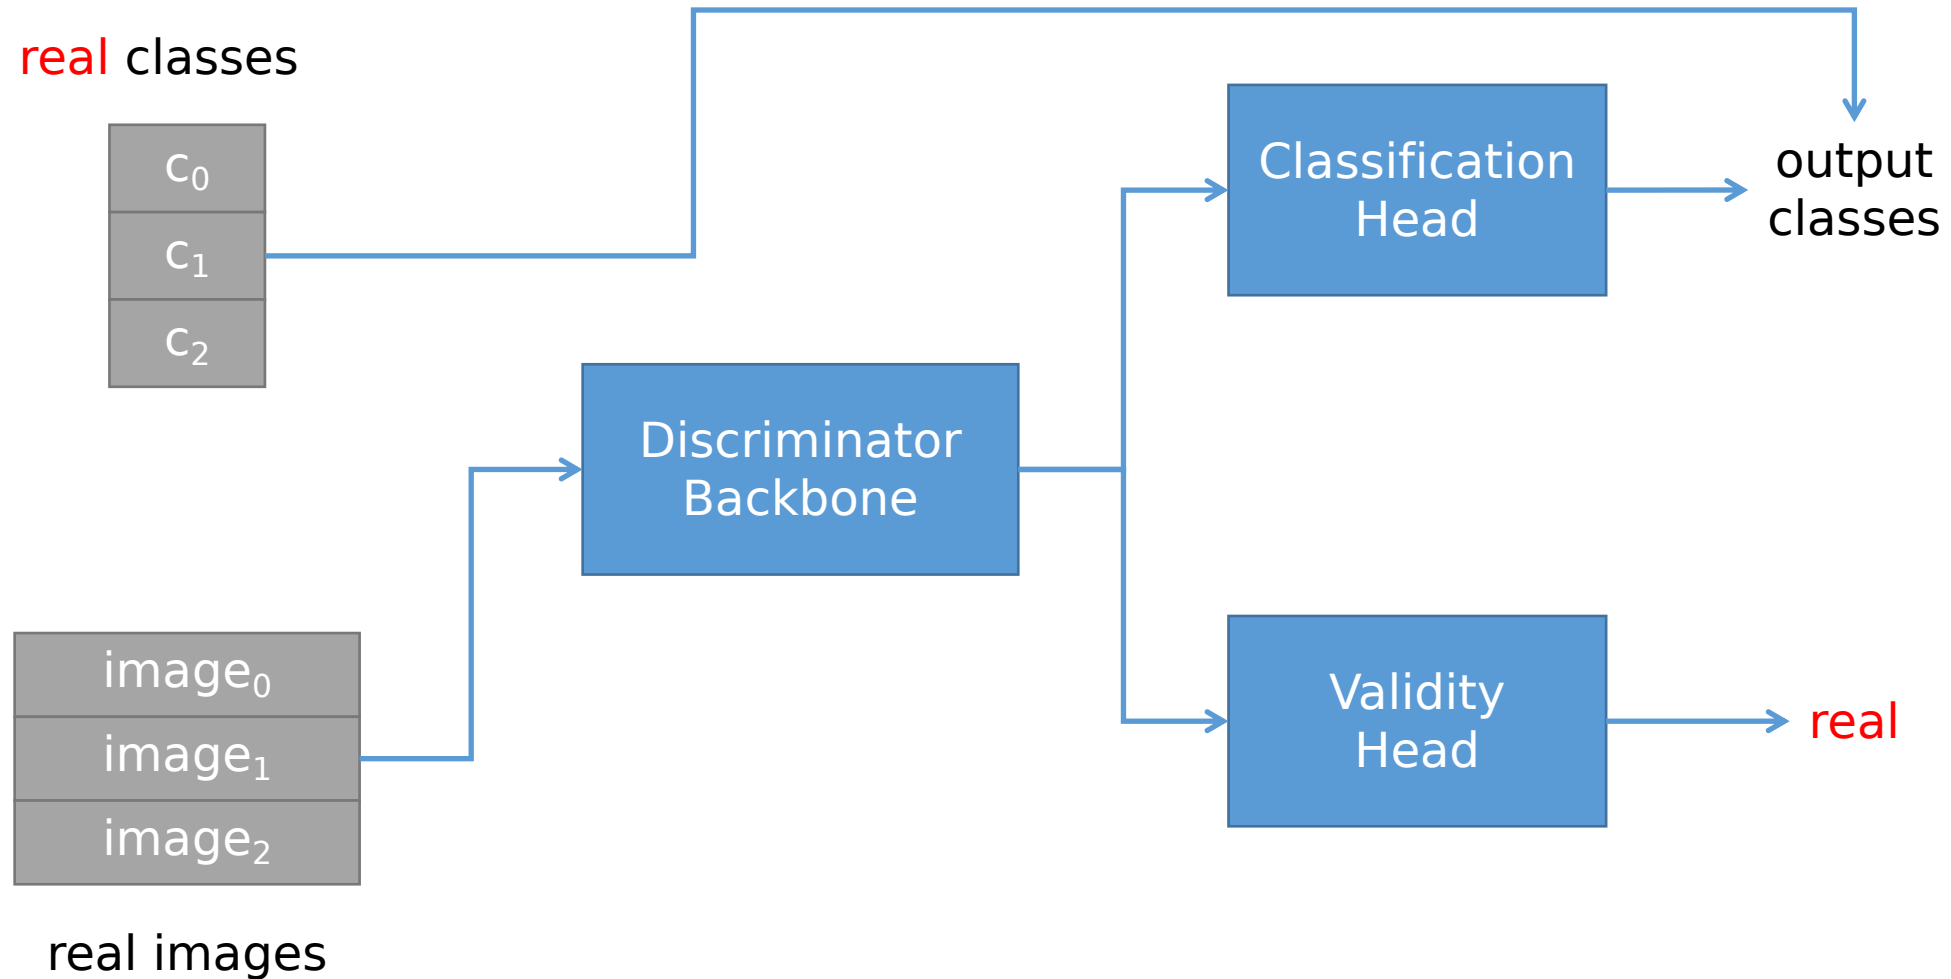


random
classes

fake



ACGANs: Discriminator Training (real images)



GAN training ticks

Training a GAN is a very **unstable** process, and **finding** the right **balance** between the **generator** and **discriminator** can be **quite tricky**.

There are some known **tricks** to help **speed** up and **stabilize** the **training process**:

- **Normalize** images between **-1** and **1**
 - Use `torchvision.transforms.Normalize((0.5,), (0.5,))` for MNIST
- **Avoid MLP, use CNN** (called **DCGAN** Deep Convolutional GAN)
- Use `torch.nn.BatchNorm2d` every layer
- **Avoid MaxPool2d, use strided** convolutions
- **Initialize convolution weights** to almost **0** (normal mean 0 and std 0.02)
- **Initialize batchnorm weights** to almost **1** (normal mean 1 and std 0.02)

GAN training ticks: in PyTorch (Initialization)

```
def init_weights(module):  
    for m in module.modules():  
        if isinstance(m, torch.nn.Conv2d):  
            torch.nn.init.normal_(m.weight, 0, 0.02)  
            torch.nn.init.constant_(m.bias, 0)  
        if isinstance(m, torch.nn.BatchNorm2d):  
            torch.nn.init.normal_(m.weight, 1.0, 0.02)  
            torch.nn.init.constant_(m.bias, 0)
```

GAN training ticks: in PyTorch (Conv block)

```
class ConvBlock(torch.nn.Module):
    def __init__(self, *, in_channels, out_channels, upsample=False, kernel=3, stride=1, padding=1):
        super(ConvBlock, self).__init__()

        self.block = torch.nn.Sequential(
            torch.nn.Upsample(scale_factor=2) if upsample else torch.nn.Identity(),
            torch.nn.Conv2d(
                in_channels,
                out_channels,
                kernel_size=kernel,
                stride=stride,
                padding=padding
            ),
            torch.nn.BatchNorm2d(out_channels),
            torch.nn.LeakyReLU()
        )

        init_weights(self.block)

    def forward(self, x):
        return self.block(x)
```

GAN training ticks: in PyTorch (Generator example)

```
class Generator(torch.nn.Module):
    def __init__(self, latent_dim, output_dim):
        super(Generator, self).__init__()
        self.latent_dim = latent_dim
        self.init_size = output_dim // 4
        self.latent_channels = 128


        self.proj_in_cnn = torch.nn.Linear(latent_dim, self.latent_channels*self.init_size**2)
        self.generator = torch.nn.Sequential(
            torch.nn.BatchNorm2d(self.latent_channels),
            ConvBlock(
                in_channels=self.latent_channels,
                out_channels=self.latent_channels//2,
                upsample=True
            ),
            ConvBlock(
                in_channels=self.latent_channels//2,
                out_channels=self.latent_channels//4,
                upsample=True
            ),
            torch.nn.Conv2d(self.latent_channels//4, 1, 3, stride=1, padding=1),
            torch.nn.Tanh()
        )

        init_weights(self.generator)
```

Exercise 1 (doable)

Train a **DCGAN** on the **MNIST** dataset

- Train the **DCGAN** for **20 epochs**
- After **each epoch** plot **10 random images generated** by the **generator**
- Tip: be gentle, use a small learning rate (like 0.0002)



Exercise 2 (very hard)

Train a **ACGAN** on the **MNIST** dataset

- Train the **ACGAN** for **20 epochs**
- After **each epoch** plot **one** random **image** for **each class generated** by the **generator**
- **Be creative**, come up with a **teqnique** to **mix** the **class** and **noise** in a way that **preserve** both **randomness** and **class information**
- Tip: If your **ACGAN** does **not** generate **conditional** images (**ignore** the **class** information) you can **weight** the **validity** and **classification losses** (0.01 for the validity loss is a good number)
- Tip: You can **use** the **torch.nn.Embedding** to **project integers** into **vectors** (like the VQ-VAE codebook)

